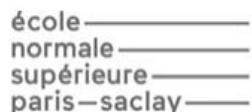


# Deep learning for medical imaging

**Olivier Colliot, PhD**  
**Research Director at CNRS**  
Co-Head of the ARAMIS Lab –  
[www.aramislab.fr](http://www.aramislab.fr)  
PRAIRIE – Paris Artificial Intelligence  
Research Institute

**Maria Vakalopoulou, PhD**  
**Assistant Professor at**  
**CentraleSupélec**  
Mathematics and Informatics (MICS)  
Office: Bouygues Building Sb.132



## Master 2 - MVA

Course website: <http://www.aramislab.fr/teaching/DLMI-2019-2020/>  
Piazza (for registered students):  
<https://piazza.com/centralesupelec/spring2020/mvadlmi/>

# Acknowledgements

---

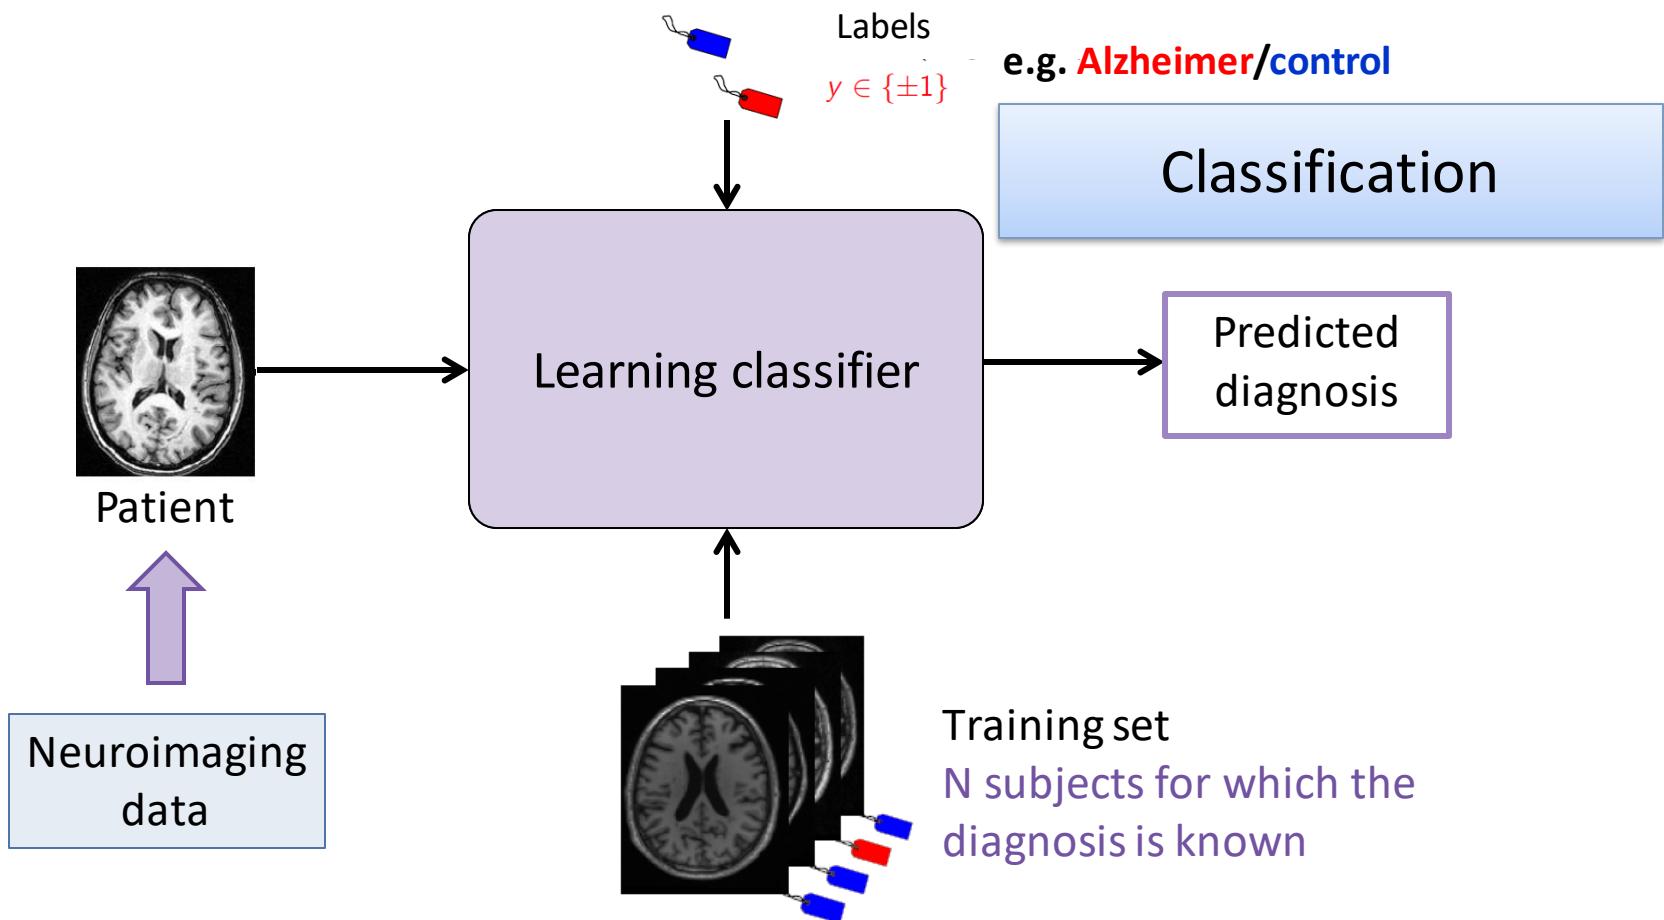
- The lecture is partially base on material by:
  - Andrej Karpathy
  - Fei-Fei
  - Daniel Rueckert
  - Ben Glocker
  - Shubhendu Trivedi & Risi Kondor

Thank you!!

# **Previous Lecture**

# **Classification & Detection**

# Example in brain image classification

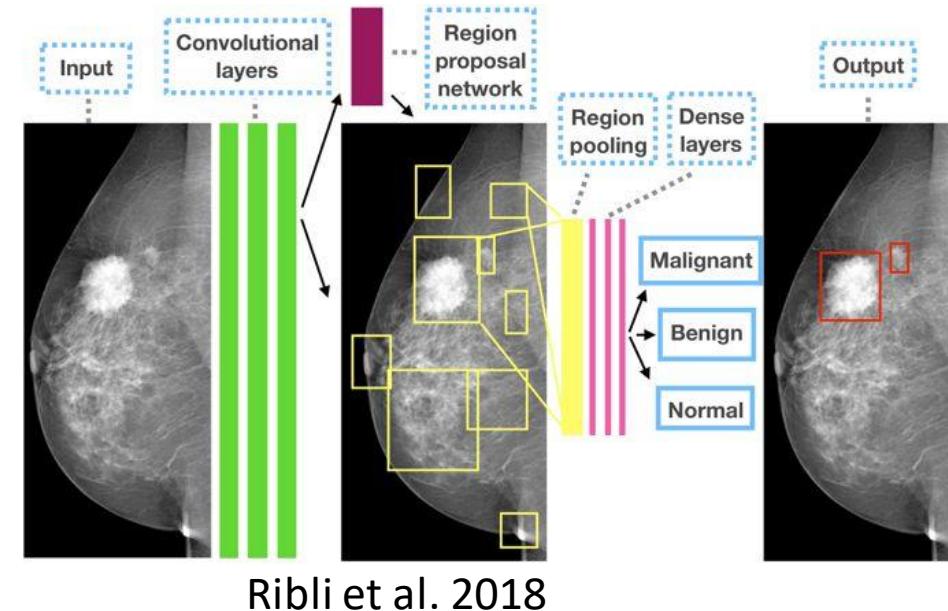


# Detecting lesions in mammograms

- Classification & Localization
- Object Detection
- R-CNN
- Fast R-CNN
- Faster R-CNN
- ....

## *Medical Imaging*

- 3D extensions
- 2 stage processes to remove false positives
- ...



# Part 4 – Segmentation

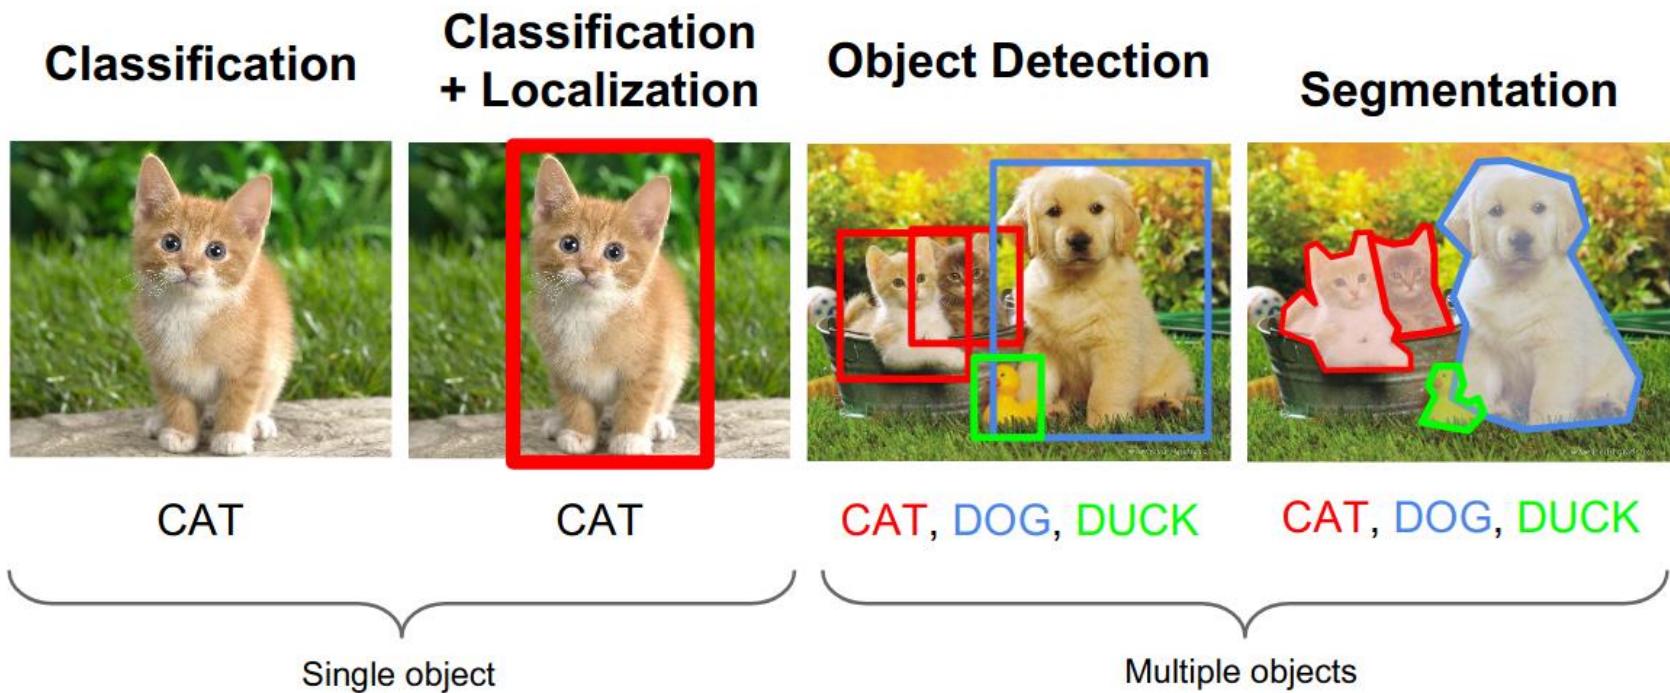
# Outline

---

- Segmentation
  - Upsampling & Dilated Filters
  - Loss Functions
  - Evaluation Metrics
- Medical Imaging

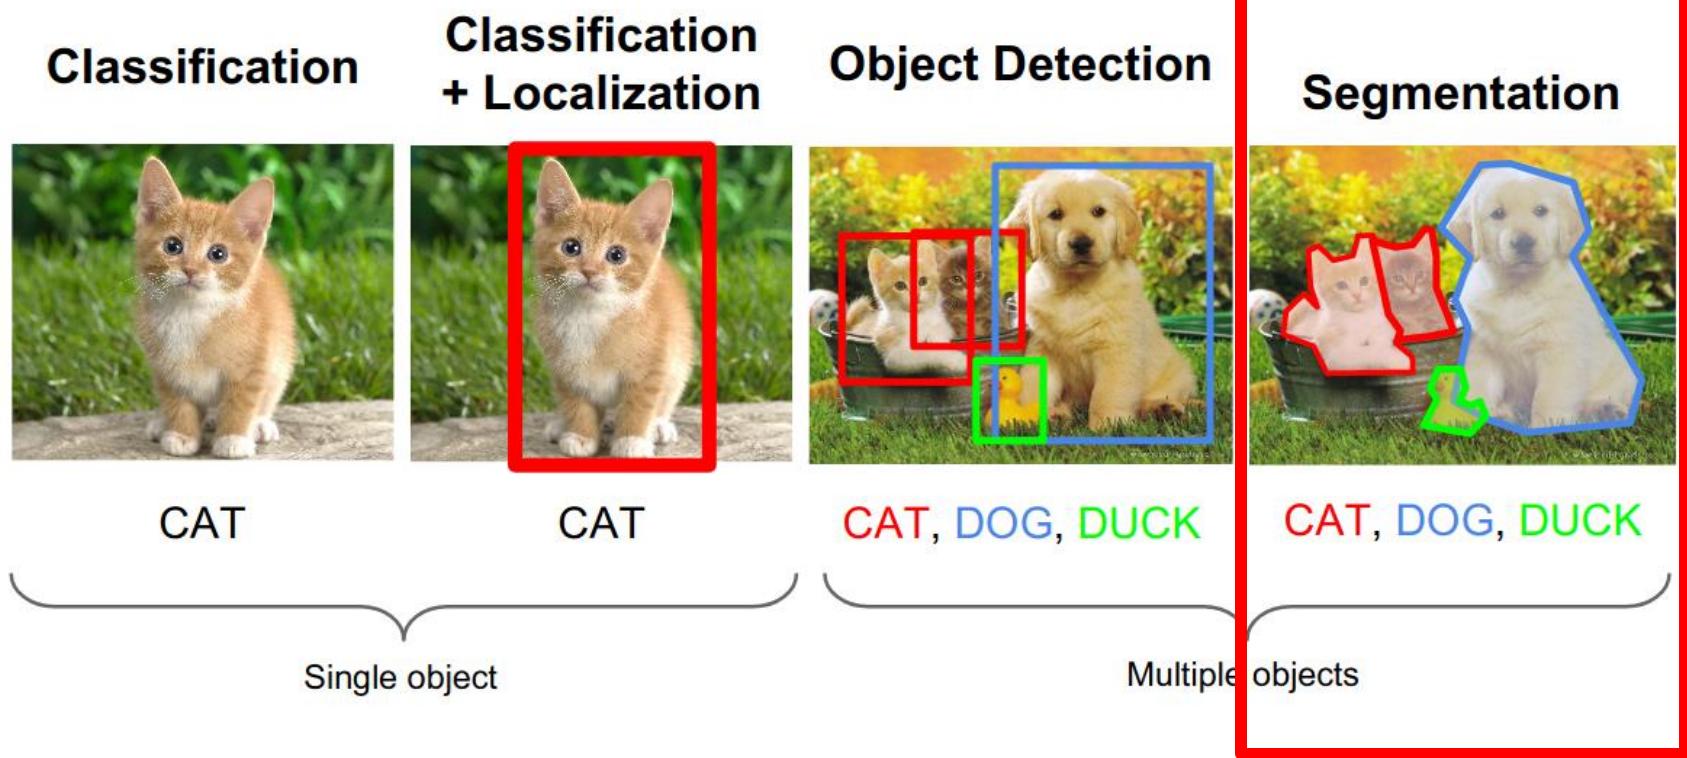
# Introduction

- Different problems for vision



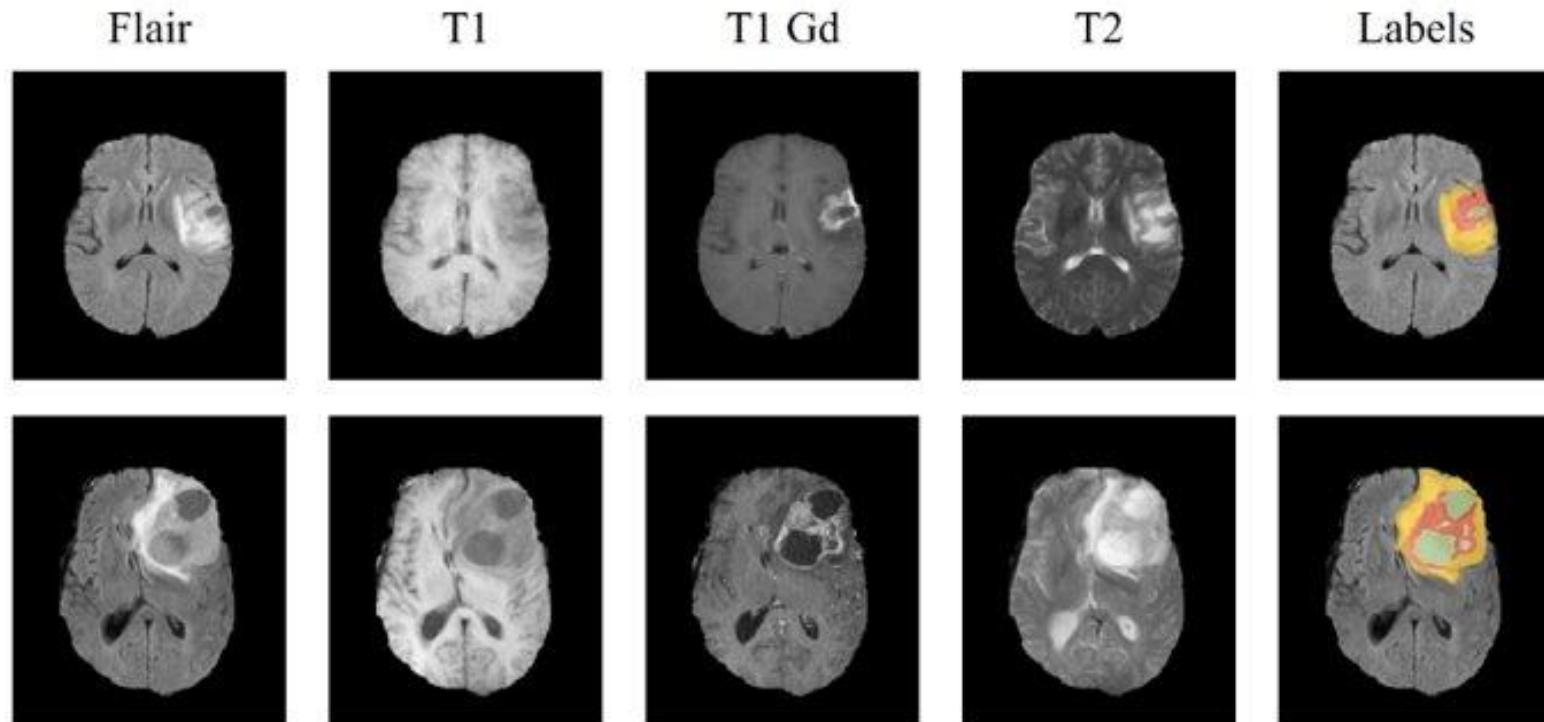
# Introduction

- Different problems for vision



# Introduction

- Semantic Segmentation
  - Label each pixel in the image with a category label
  - Do not differentiate instances only care about pixels, we want to find the mesh of pixels with the same label

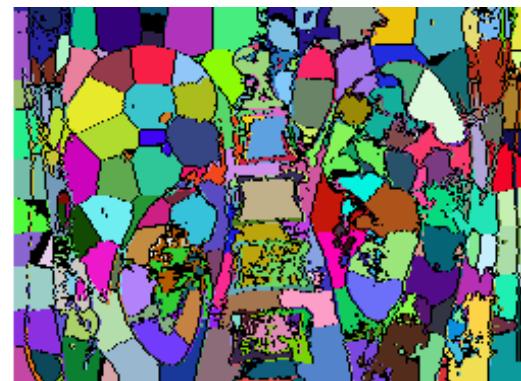


# Introduction

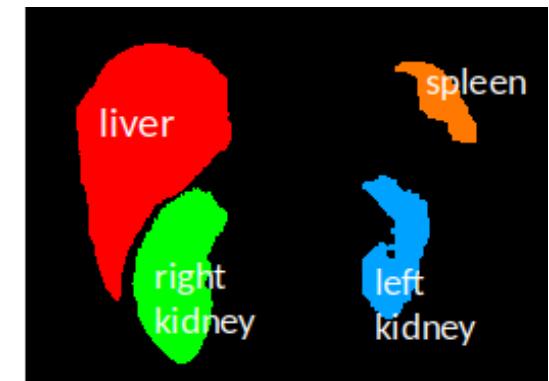
- In Semantic Segmentation, a segmented region is also assigned a **semantic meaning**
- This is in contrast to segmentation based on 'pure' clustering of image into coherent regions
- For this lecture, we are mostly interested in semantic segmentation



raw image



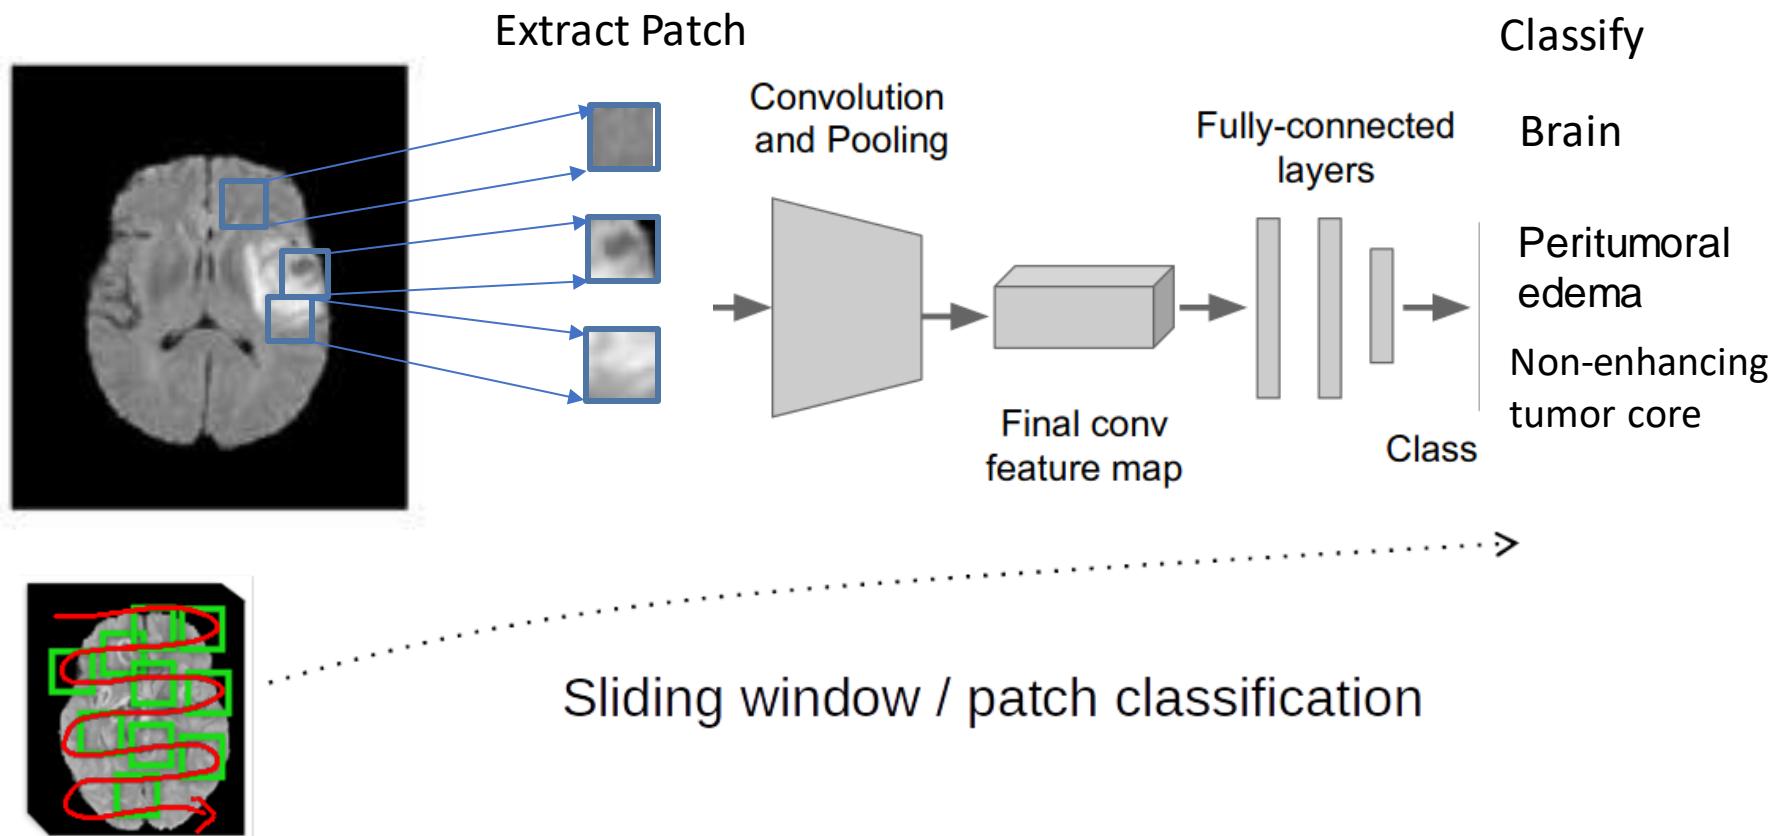
segmentation/clustering



semantic segmentation

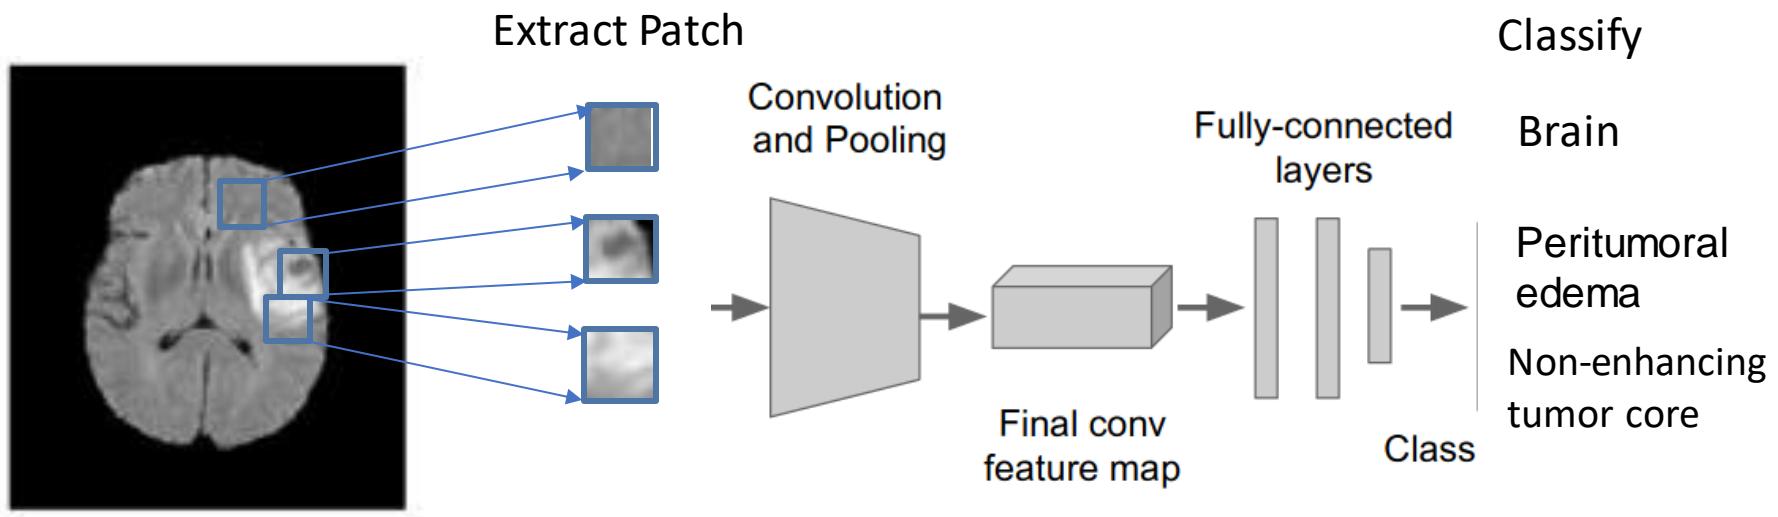
# Introduction

- Semantic Segmentation
- Idea #1: Image Segmentation as Classification using sliding window



# Introduction

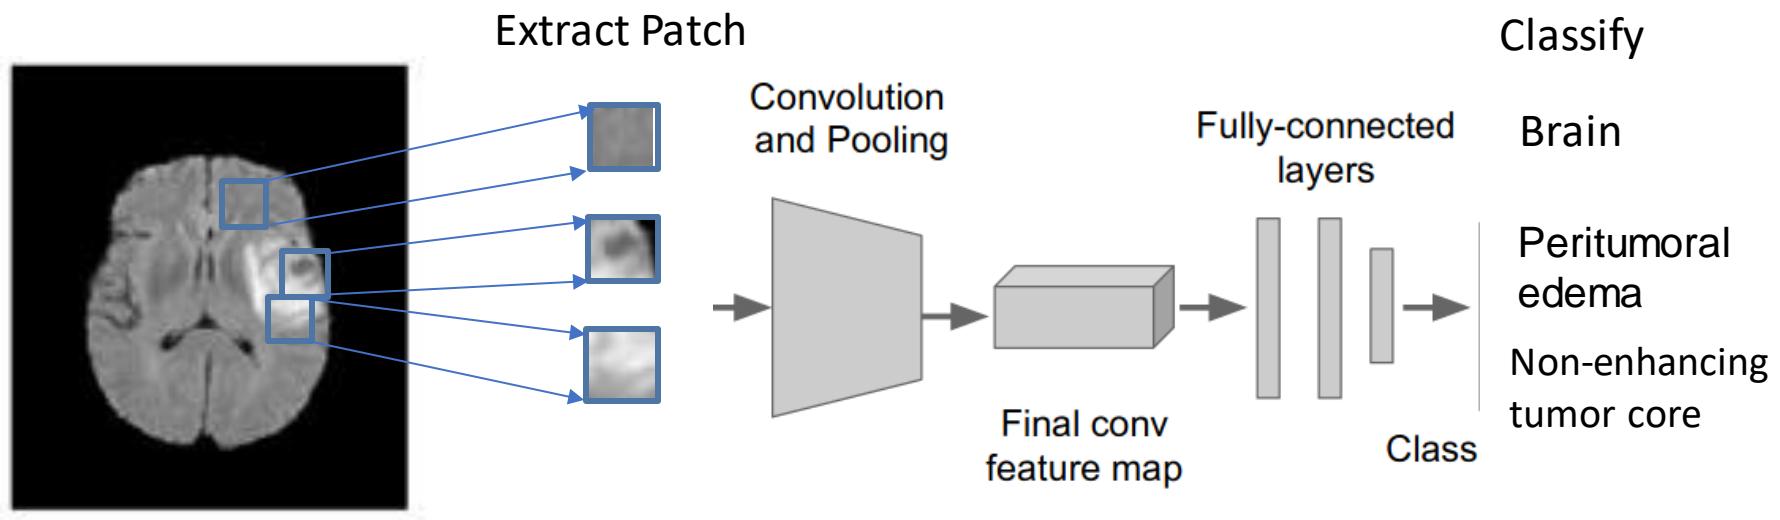
- Semantic Segmentation
- Idea #1: Image Segmentation as Classification using sliding window



- Problems??

# Introduction

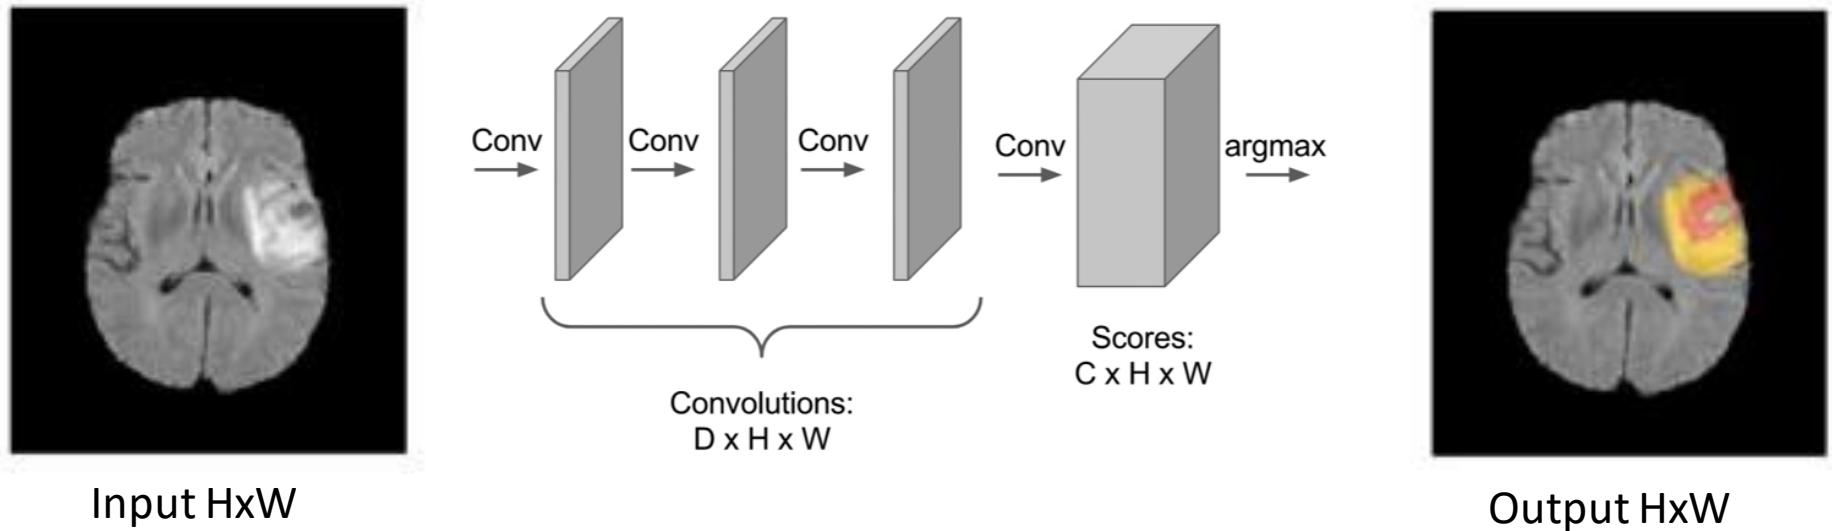
- Semantic Segmentation
- Idea #1: Image Segmentation as Classification using sliding window



- Problems??
  - Very inefficient (both time and complexity)
  - Using information only in a small part of the image
  - Do not reuse shared features between overlapping patches

# Introduction

- Semantic Segmentation
- Idea #2: Image Segmentation using Fully Convolutional Layers
  - Design a network as a bunch of convolutional layers to make predictions for pixels all at once!



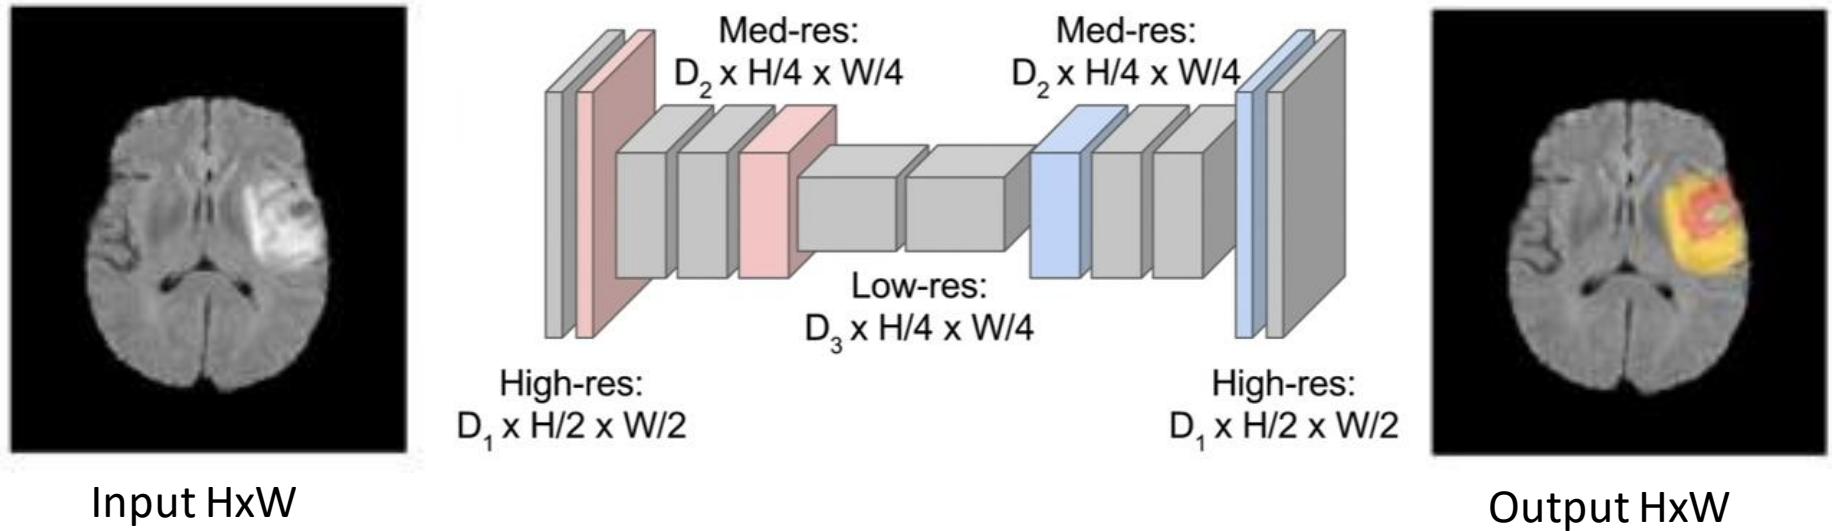
Input  $H \times W$

Output  $H \times W$

Problems?

# Introduction

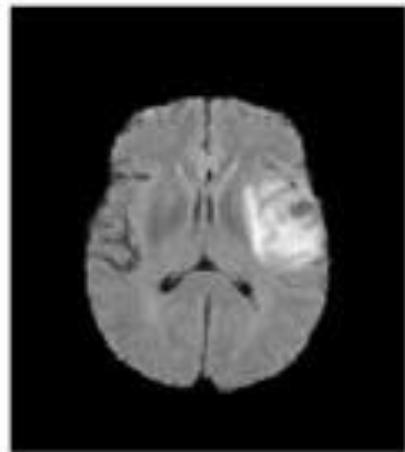
- Semantic Segmentation
- Idea #2: Image Segmentation using Fully Convolutional Layers
  - Design a network as a bunch of convolutional layers to make predictions for pixels all at once!



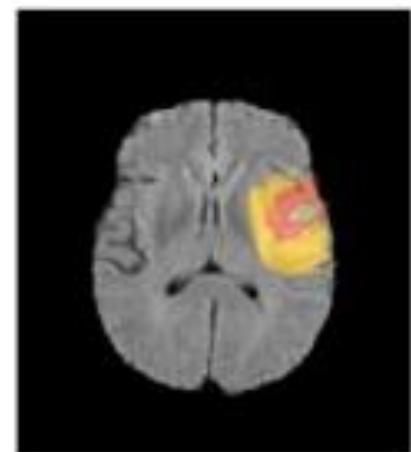
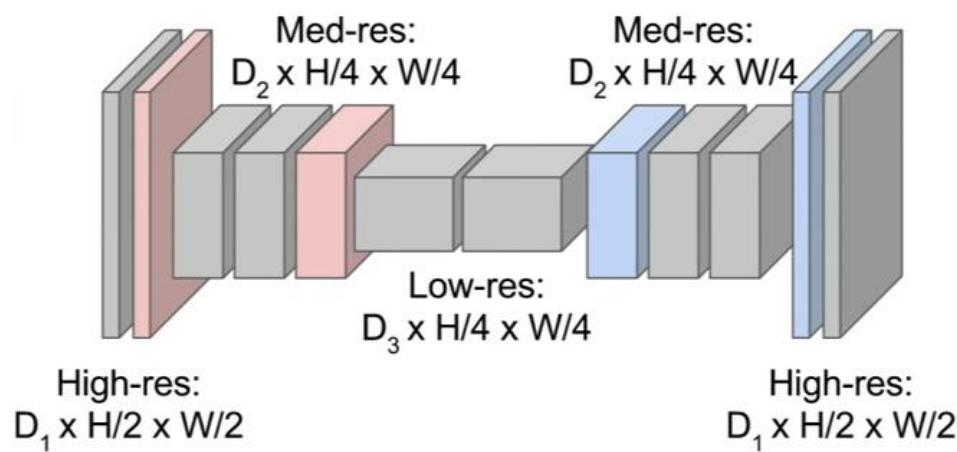
- Use Downsampling and upsampling inside the network!

# Introduction

- Semantic Segmentation
- Downsampling: Pooling, strided convolution
- Upsampling ??



Input HxW



Output HxW

# Upsampling

- In-Network upsampling: "Unpooling"

**Nearest Neighbor**

1	2
3	4

1	1	2	2
1	1	2	2
3	3	4	4
3	3	4	4

Input: 2 x 2

Output: 4 x 4

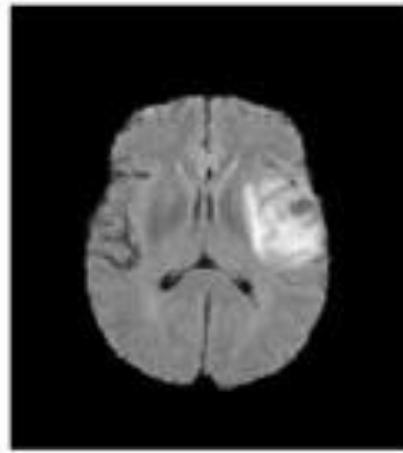
**"Bed of Nails"**

1	2
3	4

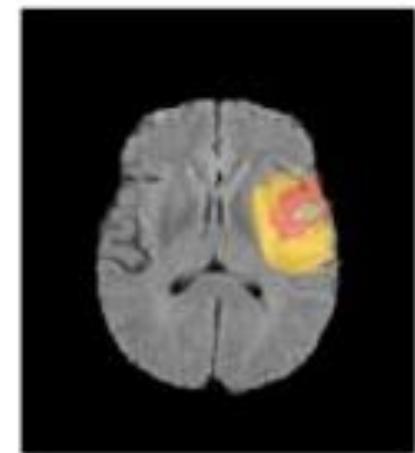
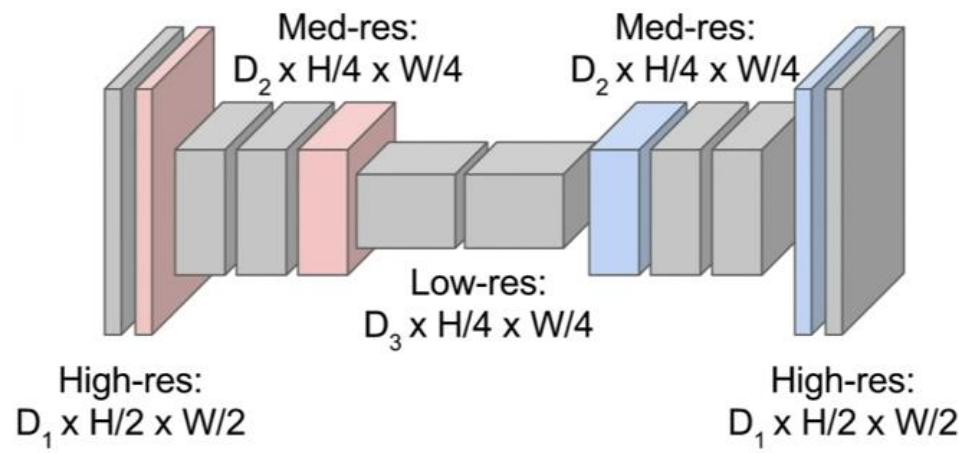
1	0	2	0
0	0	0	0
3	0	4	0
0	0	0	0

Input: 2 x 2

Output: 4 x 4



Input HxW



Output HxW

# Upsampling

- Artefacts

Transpose convolution with **zero-fill / “bed of nails”**

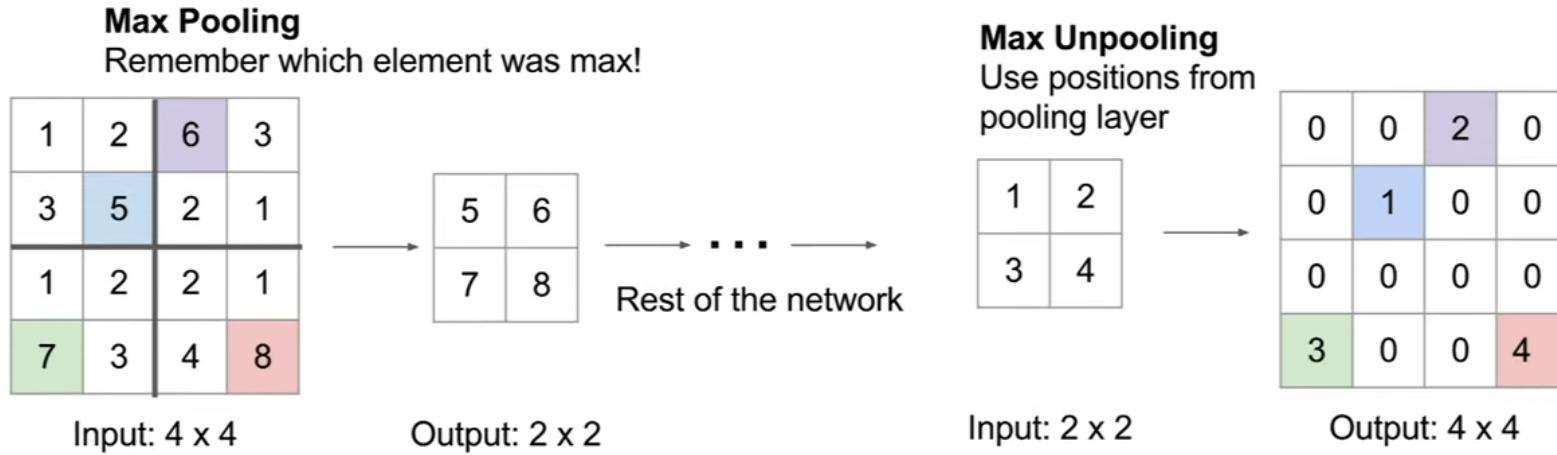


Transpose convolution with **nearest-neighbour-fill**

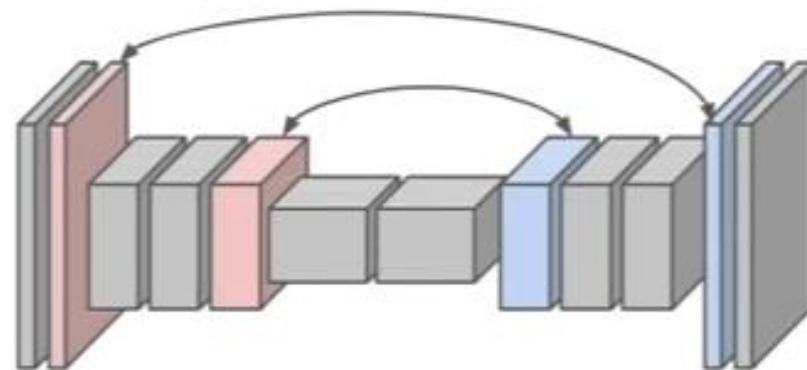


# Upsampling

- In-Network upsampling: "Max Unpooling"



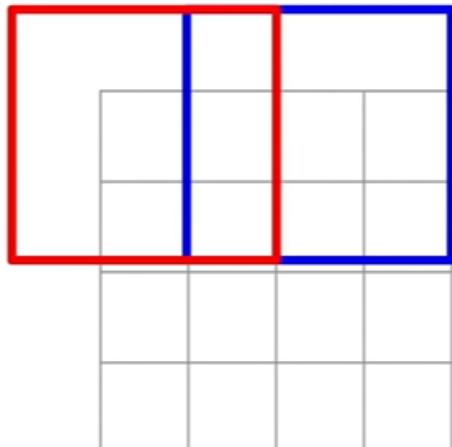
Corresponding pairs  
of downsampling and  
upsampling layers



# Upsampling

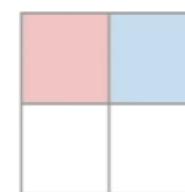
- Learnable Upsampling: Transpose Convolution (*Deconvolution*)

**Recall:** Normal  $3 \times 3$  convolution, stride 2 pad 1



Input:  $4 \times 4$

Dot product  
between filter  
and input



Output:  $2 \times 2$

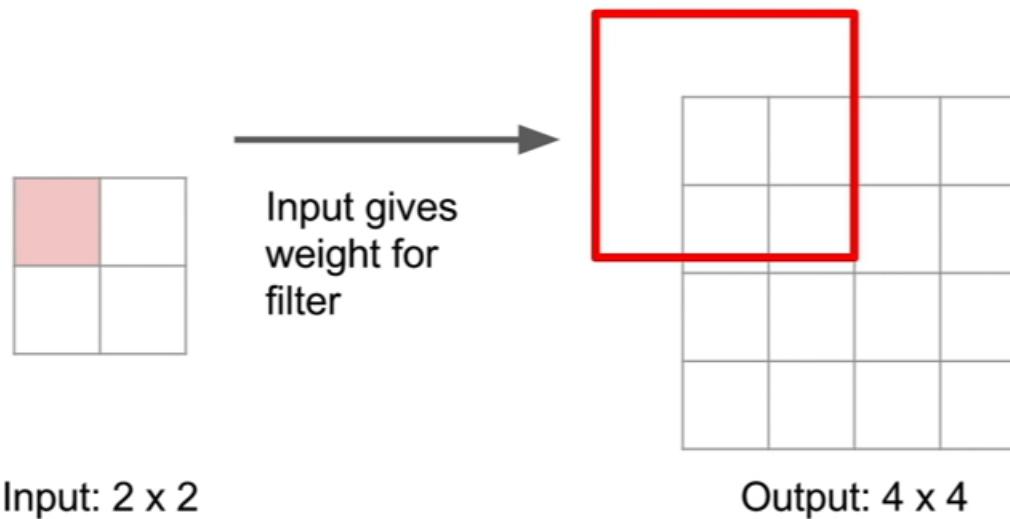
Filter moves 2 pixels in  
the input for every one  
pixel in the output

Stride gives ratio between  
movement in input and  
output

# Upsampling

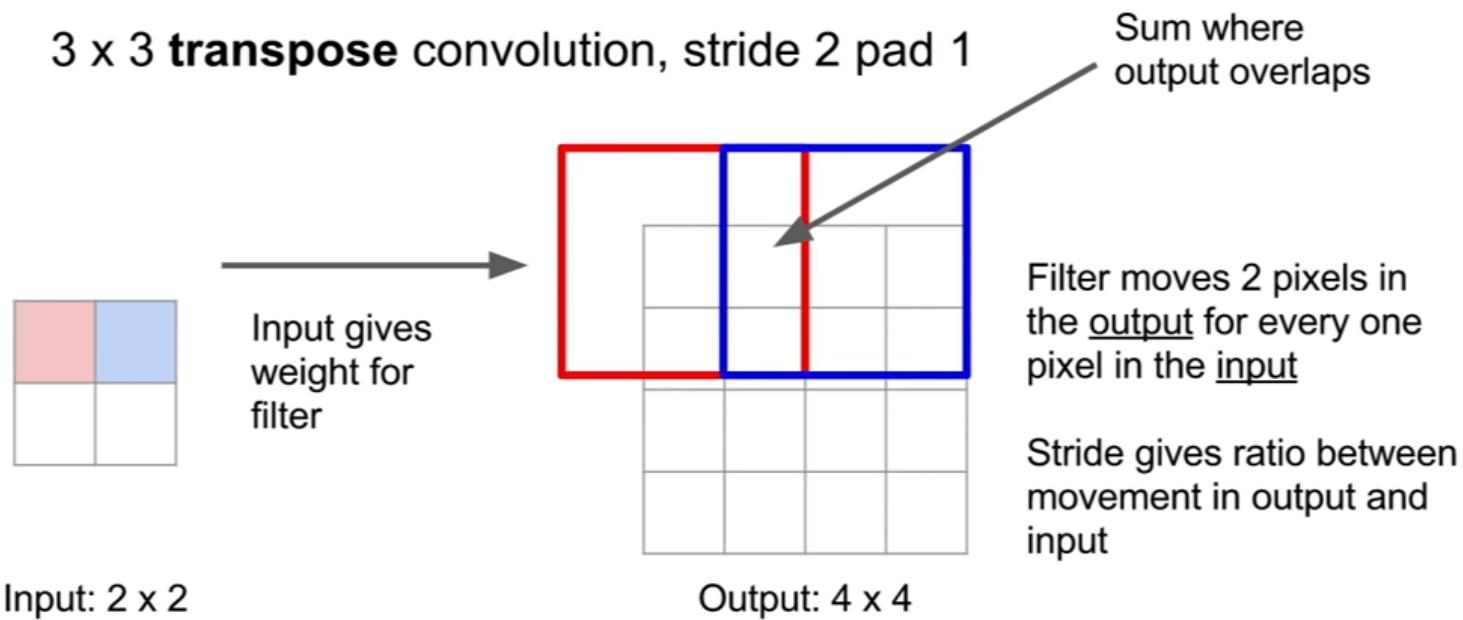
- Learnable Upsampling: Transpose Convolution (*Deconvolution*)

3 x 3 **transpose** convolution, stride 2 pad 1



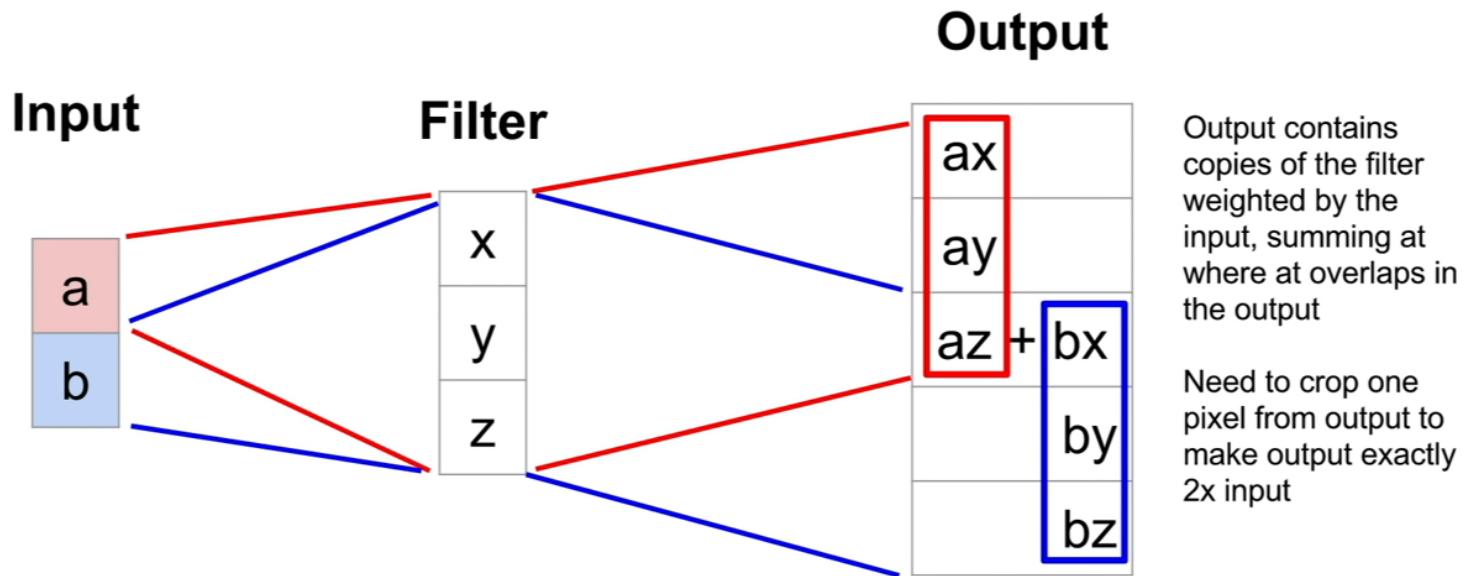
# Upsampling

- Learnable Upsampling: Transpose Convolution (*Deconvolution*)



# Upsampling

- Transpose Convolution: 1D Example



# Upsampling

- Convolution as Matrix Multiplication: 1D Example

We can express convolution in terms of a matrix multiplication

$$\vec{x} * \vec{a} = X\vec{a}$$

$$\begin{bmatrix} x & y & x & 0 & 0 & 0 \\ 0 & x & y & x & 0 & 0 \\ 0 & 0 & x & y & x & 0 \\ 0 & 0 & 0 & x & y & x \end{bmatrix} \begin{bmatrix} 0 \\ a \\ b \\ c \\ d \\ 0 \end{bmatrix} = \begin{bmatrix} ay + bz \\ ax + by + cz \\ bx + cy + dz \\ cx + dy \end{bmatrix}$$

Example: 1D conv, kernel size=3, stride=1, padding=1

# Upsampling

- Convolution as Matrix Multiplication: 1D Example

We can express convolution in terms of a matrix multiplication

$$\vec{x} * \vec{a} = X\vec{a}$$

$$\begin{bmatrix} x & y & x & 0 & 0 & 0 \\ 0 & x & y & x & 0 & 0 \\ 0 & 0 & x & y & x & 0 \\ 0 & 0 & 0 & x & y & x \end{bmatrix} \begin{bmatrix} 0 \\ a \\ b \\ c \\ d \\ 0 \end{bmatrix} = \begin{bmatrix} ay + bz \\ ax + by + cz \\ bx + cy + dz \\ cx + dy \end{bmatrix}$$

Example: 1D conv, kernel size=3, stride=1, padding=1

Convolution transpose multiplies by the transpose of the same matrix:

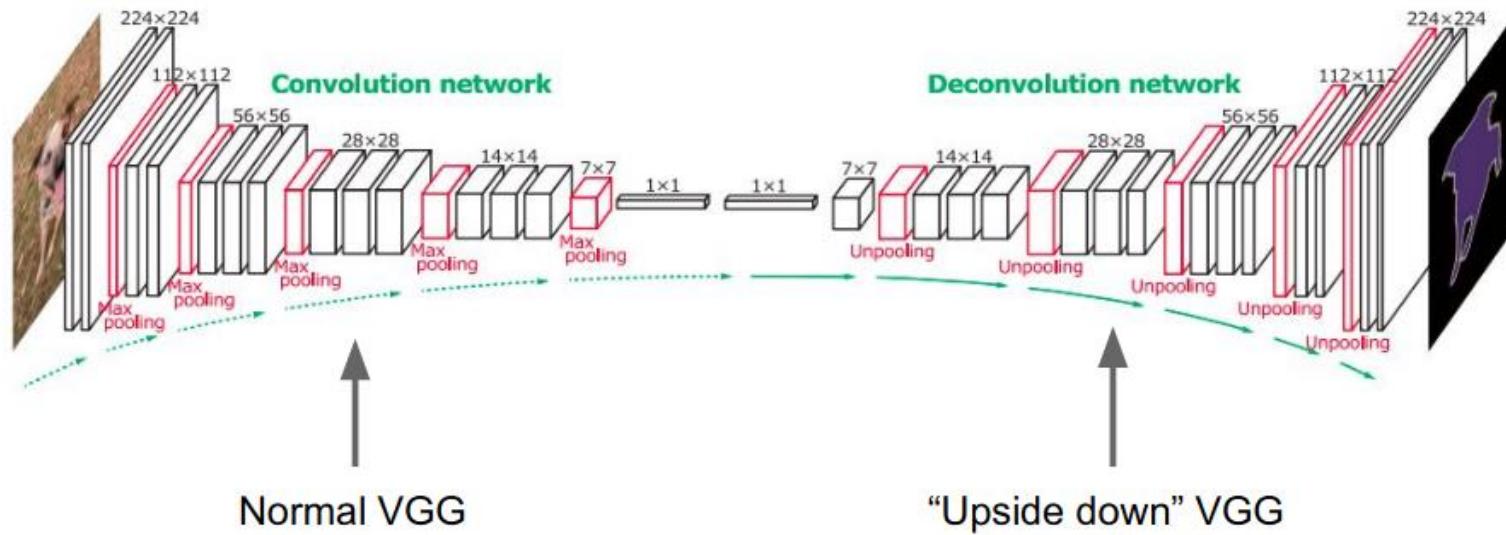
$$\vec{x} *^T \vec{a} = X^T \vec{a}$$

$$\begin{bmatrix} x & 0 & 0 & 0 \\ y & x & 0 & 0 \\ z & y & x & 0 \\ 0 & z & y & x \\ 0 & 0 & z & y \\ 0 & 0 & 0 & z \end{bmatrix} \begin{bmatrix} a \\ b \\ c \\ d \end{bmatrix} = \begin{bmatrix} ax \\ ay + bx \\ az + by + cx \\ bz + cy + dx \\ dz + dy \\ dy \end{bmatrix}$$

When stride=1, convolution transpose is just a regular convolution (with different padding rules)

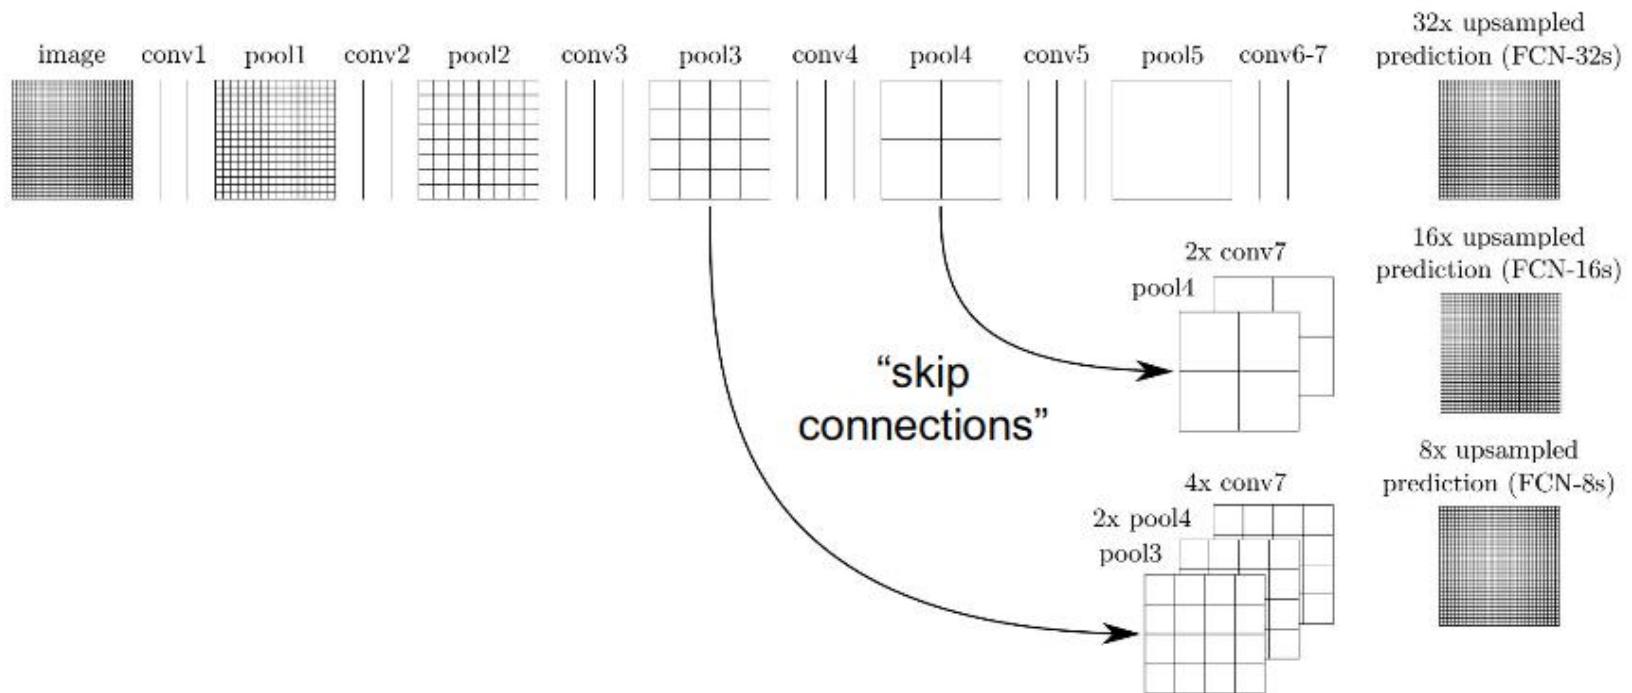
# Upsampling

- Example:



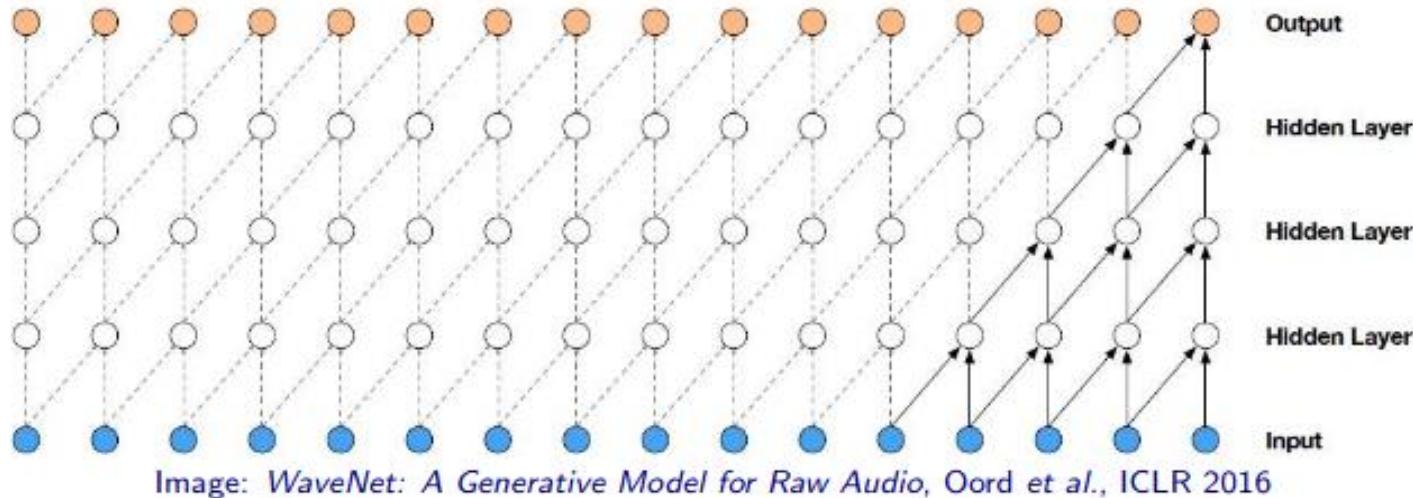
# Upsampling

- For better performance preserving more details on the boundaries usually it is combined with skip connections



# Causal Convolutions

- Connections in CNNs are sparse but units in deeper layers are connected to more of the input. At what rate does the effective receptive field size increase with depth?



- The global context of convolutional neural networks grows too slow for some applications

# UNET

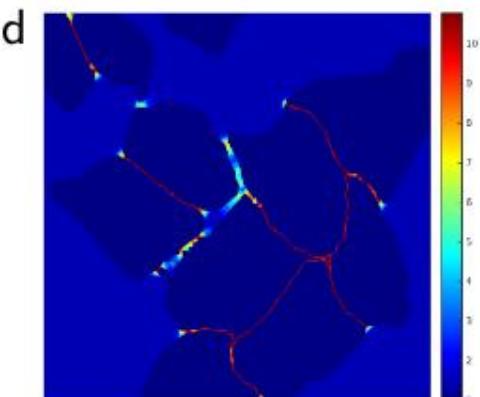
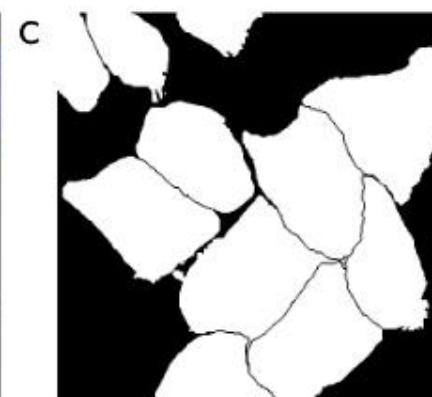
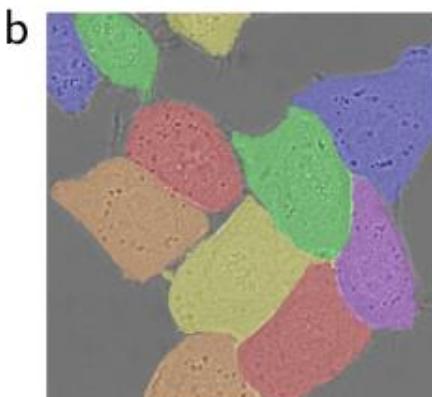
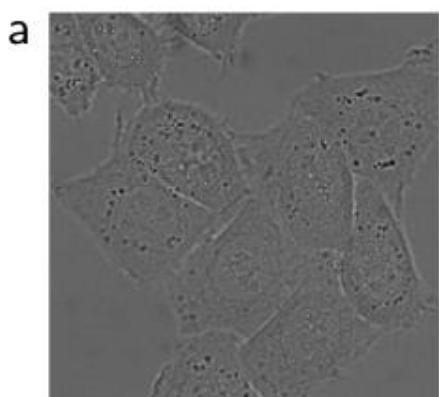
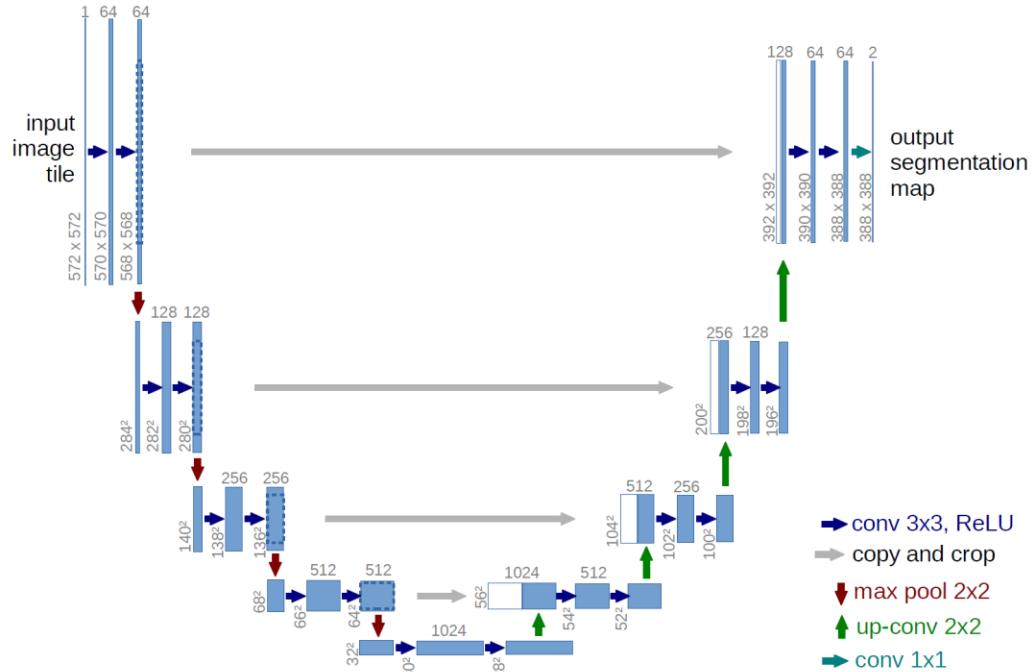
- Ronneberger et al. "U-Net: Convolutional Networks for Biomedical Image Segmentation"

## Loss function

$$E = \sum_{\mathbf{x} \in \Omega} w(\mathbf{x}) \log(p_{\ell(\mathbf{x})}(\mathbf{x}))$$

## Weighted strategy

$$w(\mathbf{x}) = w_c(\mathbf{x}) + w_0 \cdot \exp \left( -\frac{(d_1(\mathbf{x}) + d_2(\mathbf{x}))^2}{2\sigma^2} \right)$$



# 3D-UNET

- Cicek et al. "3D U-Net: Learning Dense Volumetric Segmentation from Sparse Annotation"

## Loss Function

$$E = \sum_{\mathbf{x} \in \Omega} w(\mathbf{x}) \log(p_{\ell(\mathbf{x})}(\mathbf{x}))$$

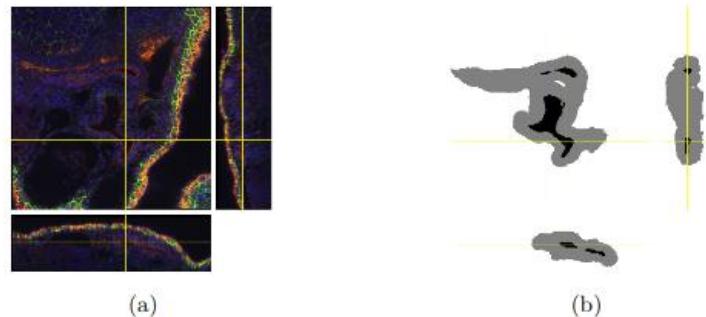


Fig. 3: (a) The confocal recording of our 3rd *Xenopus* kidney. (b) Resulting dense segmentation from the proposed 3D u-net with batch normalization.

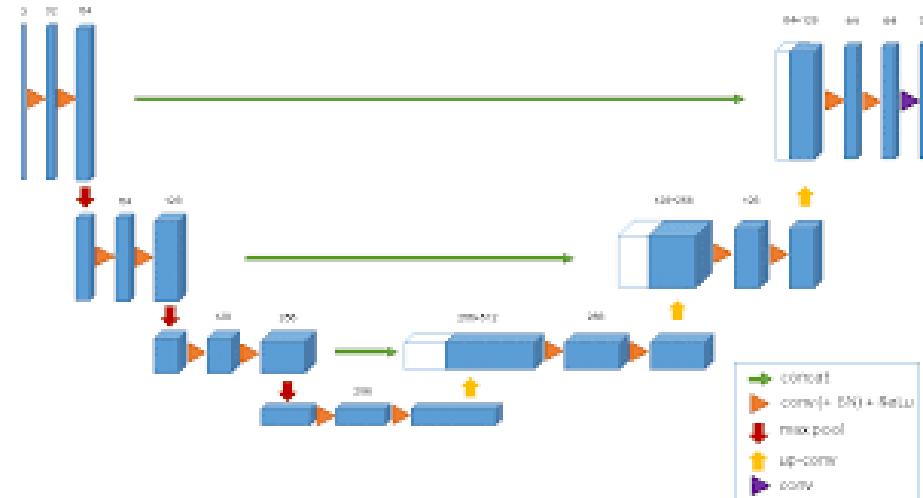


Table 1: Cross validation results for semi-automated segmentation (IoU)

test slices	3D w/o BN	3D with BN	2D with BN
subset 1	0.822	0.855	0.785
subset 2	0.857	0.871	0.820
subset 3	0.846	0.863	0.782
average	0.842	0.863	0.796

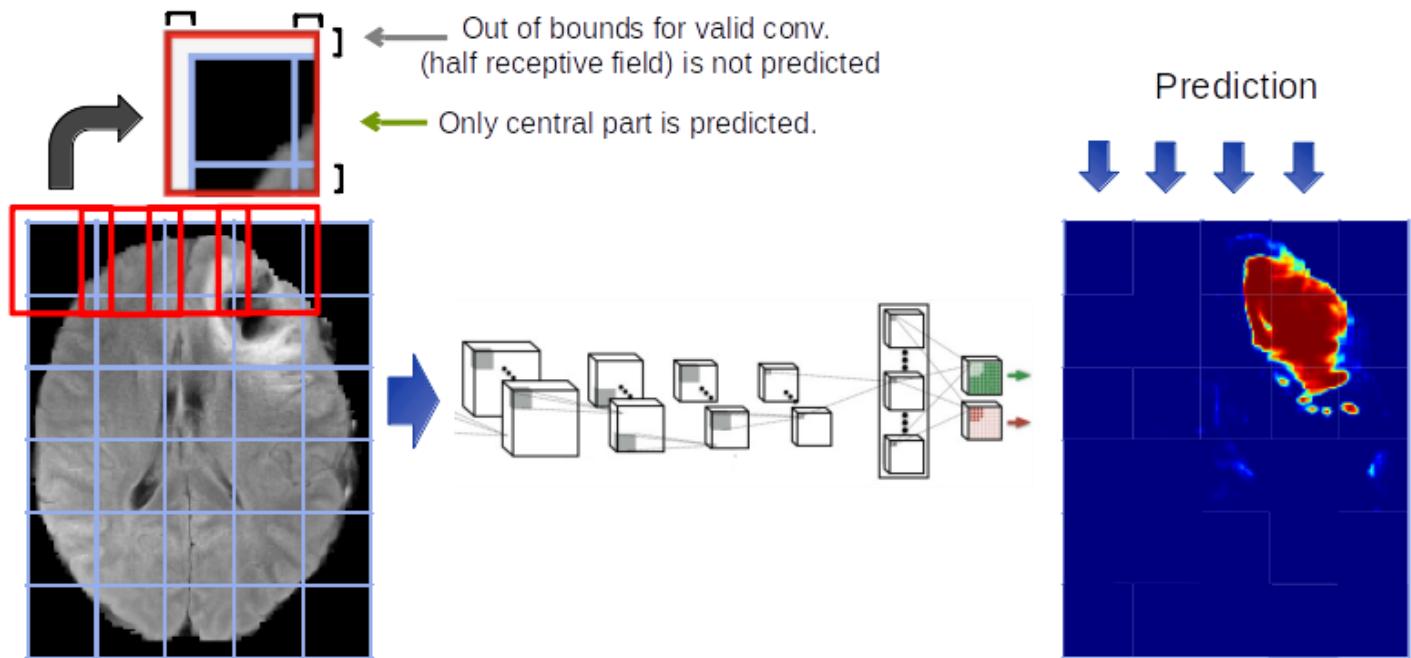
Table 2: Effect of # of slices for semi-automated segmentation (IoU)

GT slices	GT voxels	IoU S1	IoU S2	IoU S3
1,1,1	2.5%	0.331	0.483	0.475
2,2,1	3.3%	0.676	0.579	0.738
3,3,2	5.7%	0.761	0.808	0.835
5,5,3	8.9%	0.856	0.849	0.872

Table 3: Cross validation results for fully-automated segmentation (IoU)

test volume	3D w/o BN	3D with BN	2D with BN
1	0.655	0.761	0.619
2	0.734	0.798	0.698
3	0.779	0.554	0.325
average	0.723	0.704	0.547

# Tiling for 3D networks



Bigger input tiles (or segments):

- Require more memory (FMs expand)
- But fewer redundant computations (less overlap)

Tile size during testing can be different from training

# Dilated Convolutions

---

- Recall discrete convolution

$$S(i, j) = (I * K)(i, j) = \sum_m \sum_n I(m, n)K(i - m, j - n)$$

- Dilated Convolution

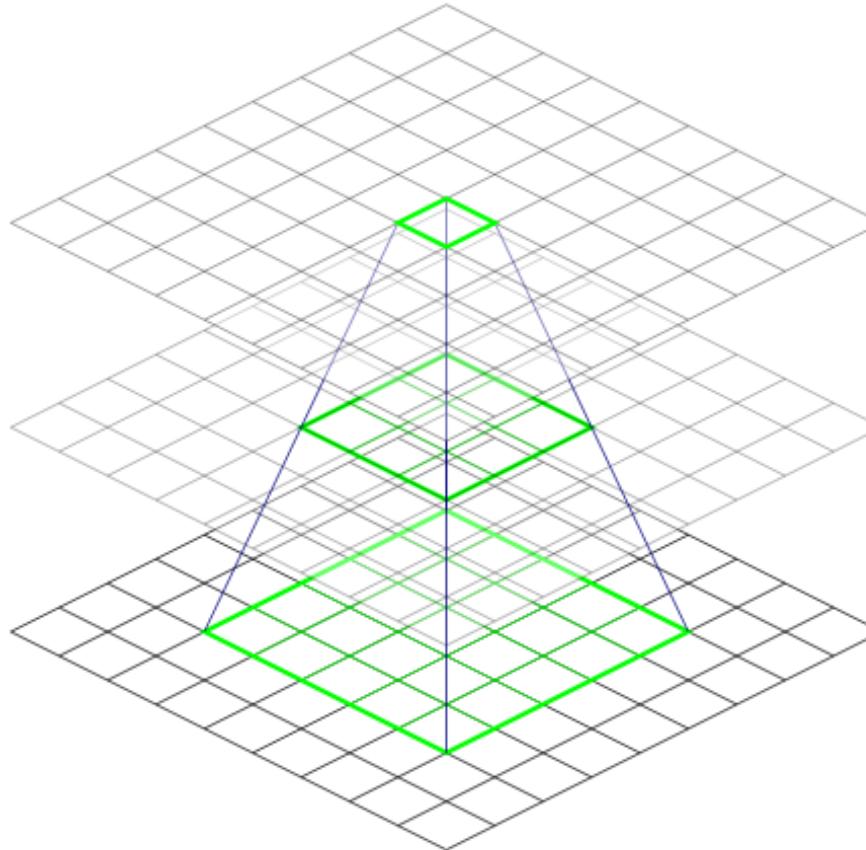
$$S(i, j) = (I *_l K)(i, j) = \sum_m \sum_n I(m, n)K(i - lm, j - ln)$$

L is a dilation factor

- Very old idea going to the 80s wavelet theory literature

# Dilated Convolutions

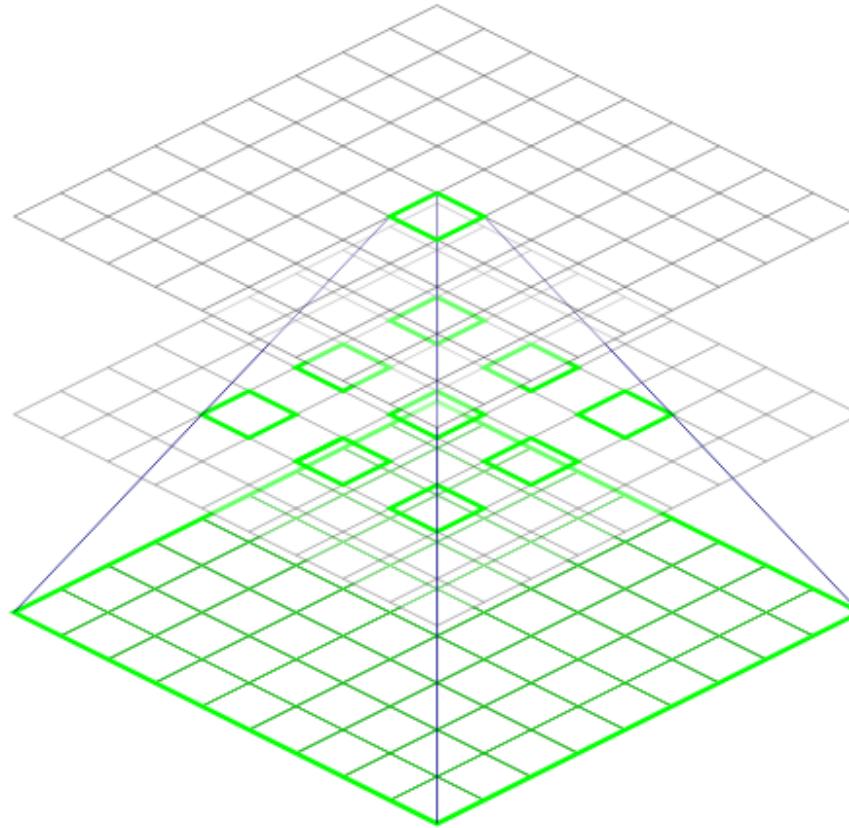
- Regular Convolution



The unit on the third layer has an effective receptive field of size 5x5

# Dilated Convolutions

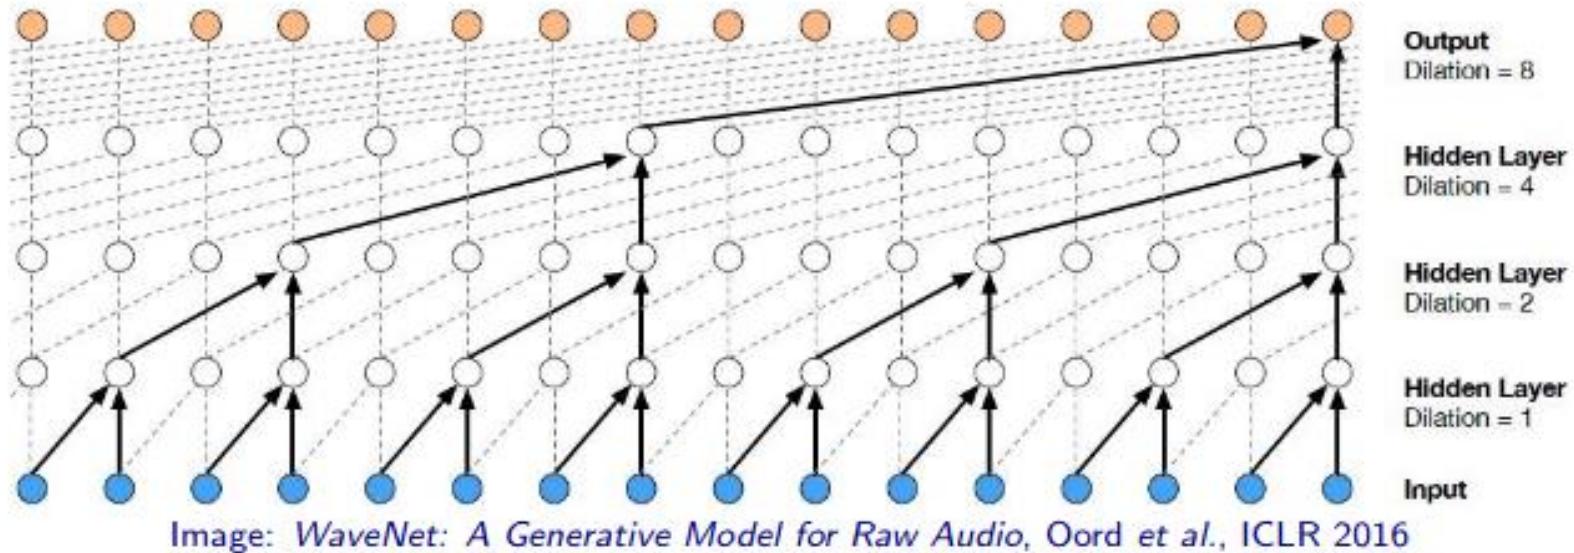
- Dilated Convolution



The unit on the third layer has an effective receptive field of size 9x9

# Dilated Convolutions

- Dilated Convolution



# HighRes3DNet

- Li et. al, "On the Compactness, Efficiency, and Representation of 3D Convolutional Networks: Brain Parcellation as a Pretext Task", IPMI 2017

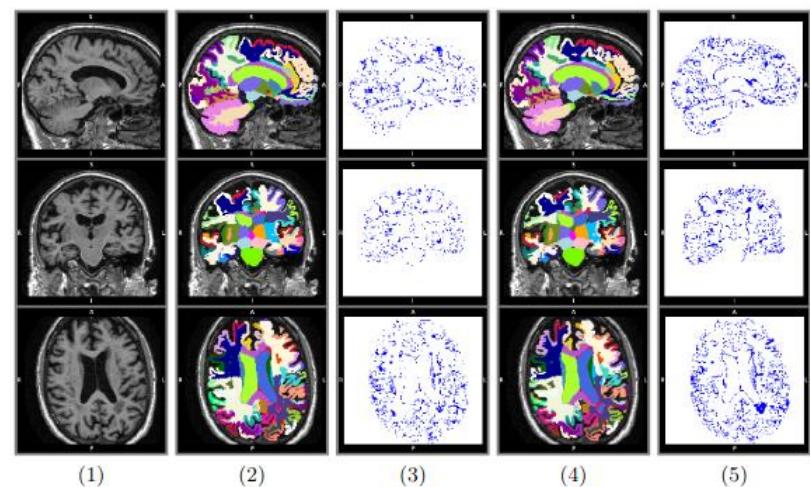
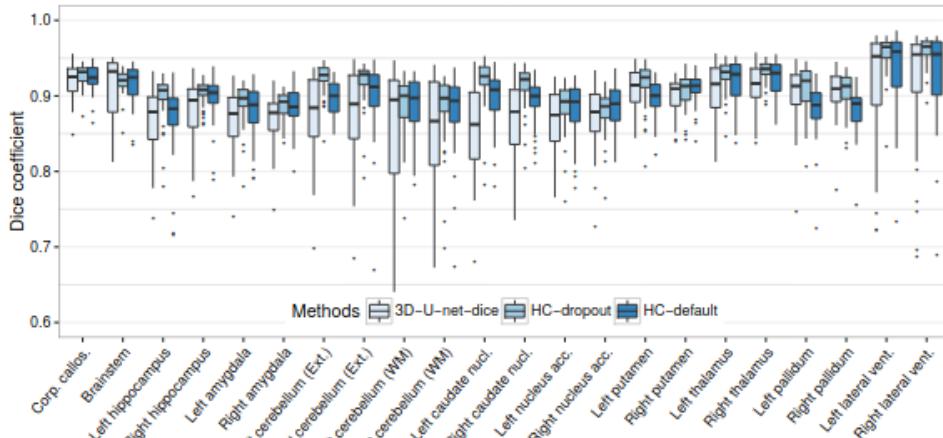
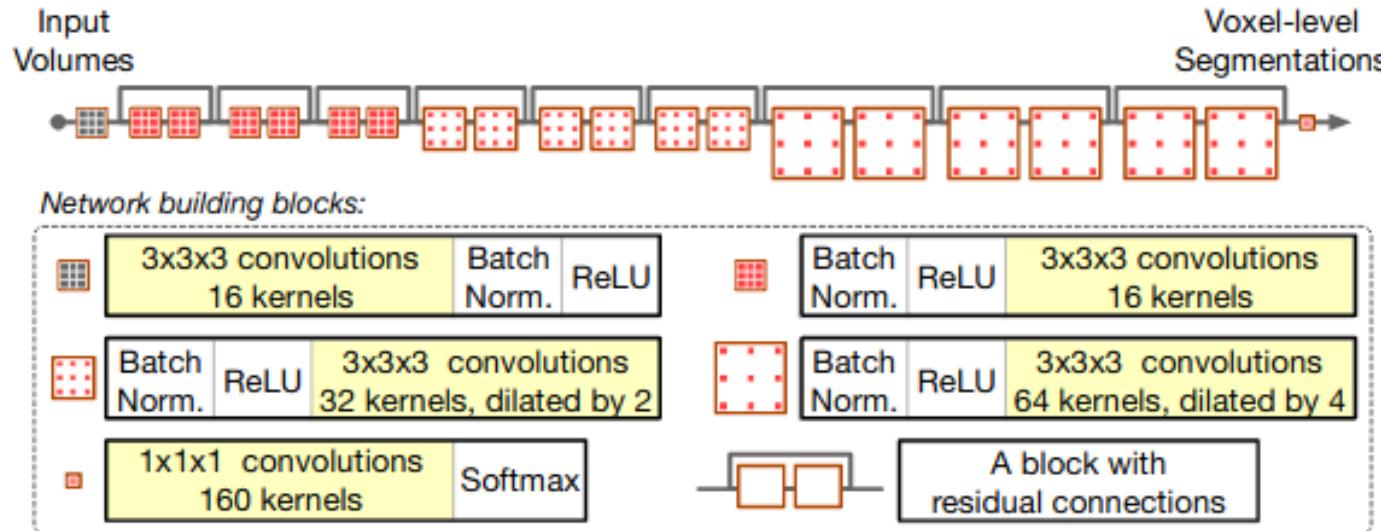
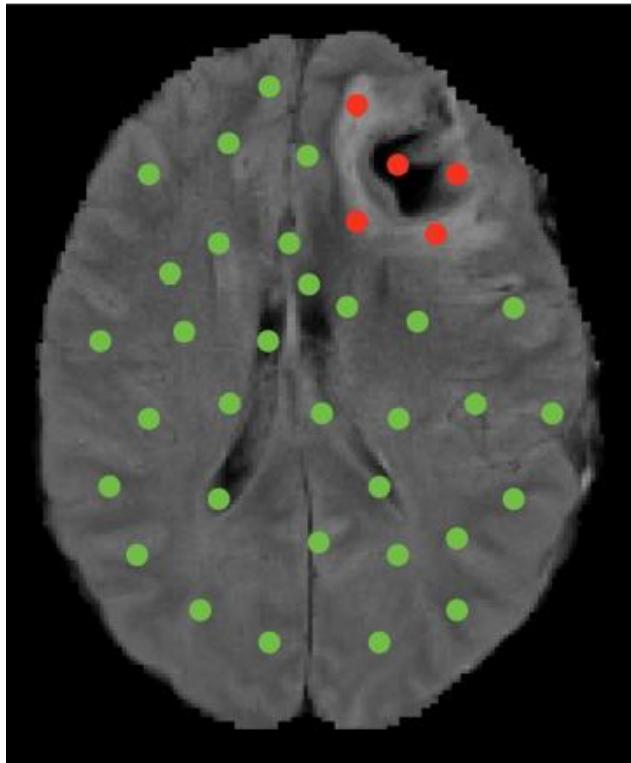


Fig. 4. Visualisations of segmentation results. (1) slices from a test image volume, segmentation maps and false prediction maps generated by HC-dropout (2, 3), and 3D U-net-dice (4, 5).

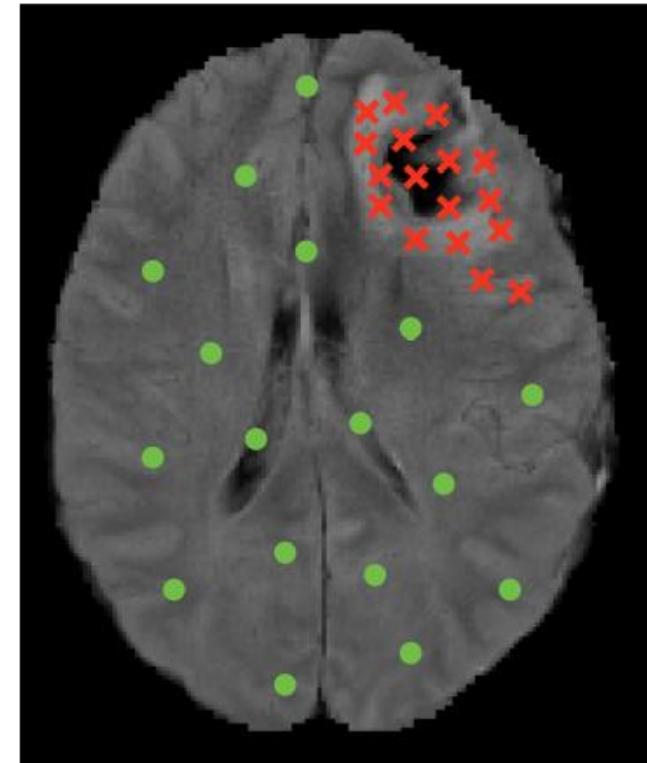
# Loss Functions

# Class Imbalance and Sampling

Uniform sampling



Weighted sampling: 50/50



# Training

## Training with single predictions on each sample:

Learn parameters  $\theta = \min_{\theta} \mathcal{L}$

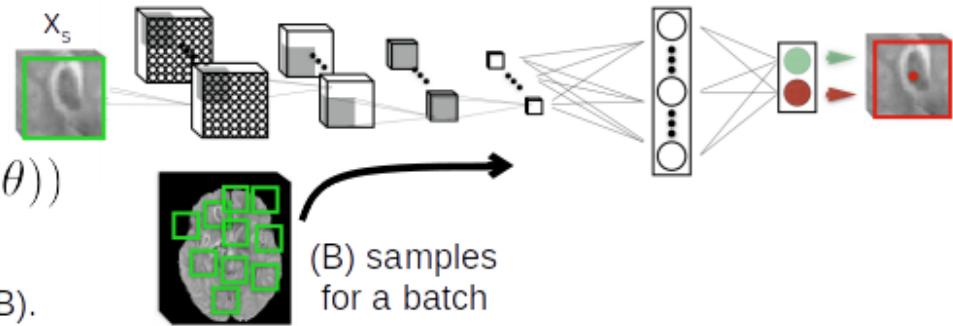
$$\mathcal{L} = -\frac{1}{B} \sum_{s=1}^B \sum_c [c = y_s] \log(f_c(x_s; \theta))$$

Cross Entropy over all the samples of a batch (B).

$x_s$  a training sample (patch),  $y_s$  the true label of its central voxel.

$f_c(x)$  the CNN's confidence for predicting class  $c$ .

[.] the indicator function.

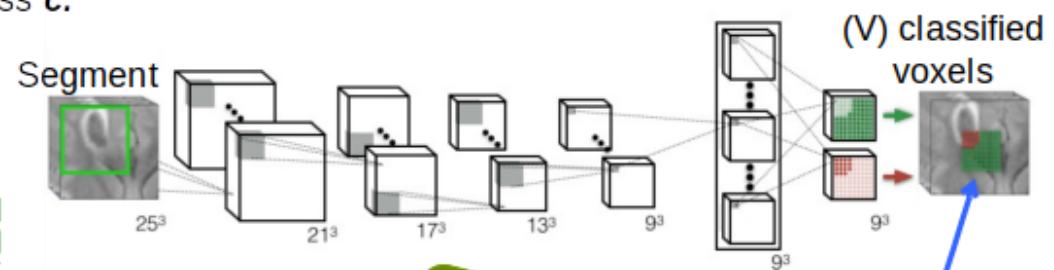


## Training on dense predictions:

Over the (B) segments (s) in a batch.

$$\mathcal{L} = -\frac{1}{B \cdot V} \sum_{s=1}^B \sum_{v=1}^V \sum_c [c = y_s^v] \log(f_c(x_s^v; \theta))$$

Over the (V) classified voxels in a segment.



Increase number of samples per batch.  
Without linear increase in computation  
(eg if loading V separate patches).

# Additional Losses

---

- **Dice Loss [Milletari et al., 2016]**

$$D = \frac{2 \sum_i^N p_i g_i}{\sum_i^N p_i^2 + \sum_i^N g_i^2}$$
$$\frac{\partial D}{\partial p_j} = 2 \left[ \frac{g_j \left( \sum_i^N p_i^2 + \sum_i^N g_i^2 \right) - 2p_j \left( \sum_i^N p_i g_i \right)}{\left( \sum_i^N p_i^2 + \sum_i^N g_i^2 \right)^2} \right]$$

$g_n$  the value of segmentation at n and  $p_n$  the predicted probabilistic map

# Additional Losses

---

- **Dice Loss** [Milletari et al., 2016]

$$D = \frac{2 \sum_i^N p_i g_i}{\sum_i^N p_i^2 + \sum_i^N g_i^2} \quad \frac{\partial D}{\partial p_j} = 2 \left[ \frac{g_j \left( \sum_i^N p_i^2 + \sum_i^N g_i^2 \right) - 2p_j \left( \sum_i^N p_i g_i \right)}{\left( \sum_i^N p_i^2 + \sum_i^N g_i^2 \right)^2} \right]$$

$g_n$  the value of segmentation at n and  $p_n$  the predicted probabilistic map

- **Generalized Dice Loss** [Sudre et al.. 2017]

$$GDL = 1 - 2 \frac{\sum_{l=1}^2 w_l \sum_n r_{ln} p_{ln}}{\sum_{l=1}^2 w_l \sum_n r_{ln} + p_{ln}}, \quad \frac{\partial GDL}{\partial p_i} = -2 \frac{(w_1^2 - w_2^2) \left[ \sum_{n=1}^N p_n r_n - r_i \sum_{n=1}^N (p_n + r_n) \right] + N w_2 (w_1 + w_2) (1 - 2r_i)}{\left[ (w_1 - w_2) \sum_{n=1}^N (p_n + r_n) + 2N w_2 \right]^2}$$

$r_n$  the value of segmentation at n,  $p_n$  the predicted probabilistic map,  $w_l = 1 / \sum(r_{ln})^2$

*Used to deal with the correlation that exists between the size of the segment and the result of the dice!*

# Additional Losses

- Most of the losses balance the importance of positive/negative examples without integrating the importance between easy/hard examples.
- Focal Loss** [Lin et al. 2018]
  - Propose to reshape the loss function to down-weight easy examples and thus focus training on hard negatives.

*Binary Cross entropy*

$$\text{CE}(p, y) = \begin{cases} -\log(p) & \text{if } y = 1 \\ -\log(1-p) & \text{otherwise.} \end{cases}$$

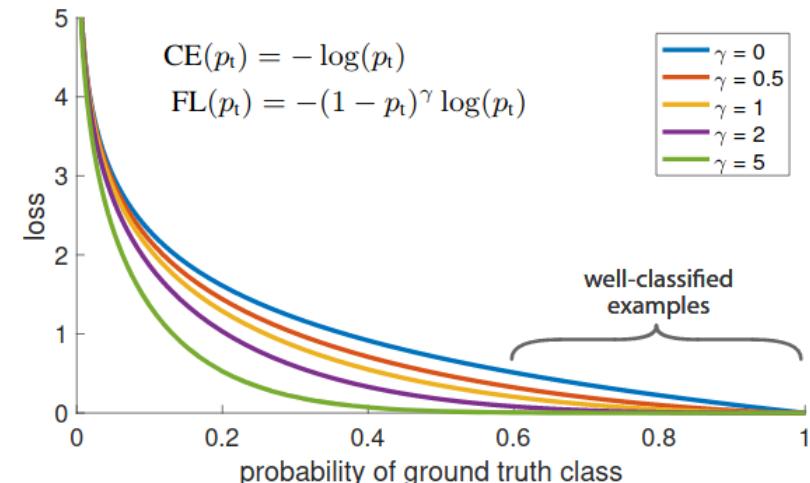
*Focal Loss*

$$\text{FL}(p_t) = -(1 - p_t)^\gamma \log(p_t).$$

Where  $\gamma \geq 0$ .

when an example is misclassified and  $p$  is small

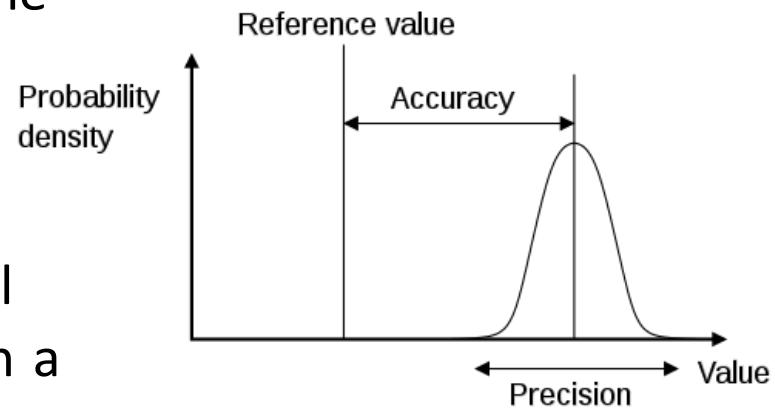
the modulating factor is near 1 and the loss is unaffected. As  $p \rightarrow 1$ , the factor goes to 0 and the loss of the well-classified examples is down-weighted



# Evaluation Metrics

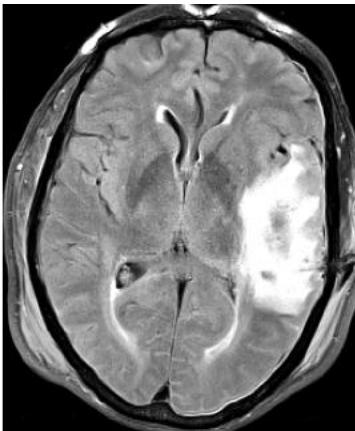
# How to Assess Performance?

- **Precision**
  - is a description of random errors, a measure of statistical variability.
  - the repeatability, or reproducibility of the measurement
- **Accuracy**
  - More commonly, it is a description of systematic errors, a measure of statistical bias; as these cause a difference between a result and a "true" value, ISO calls this trueness.
- **Robustness**
  - Refers to the degradation in performance with respect to varying noise levels or other imaging artefacts

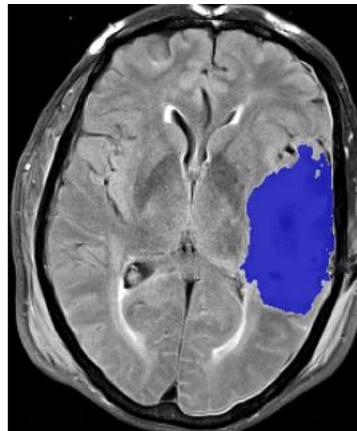


# How to Assess Performance?

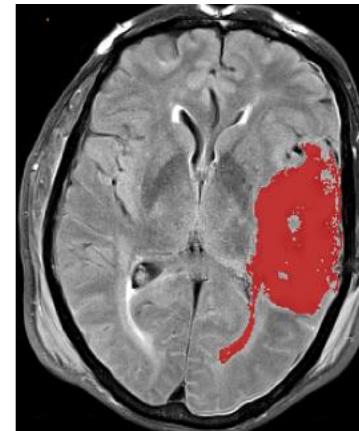
- Assume we have an automated segmentation algorithm, how do we assess "how good it is"?
- There are several ways of quantifying segmentation performance



MR image



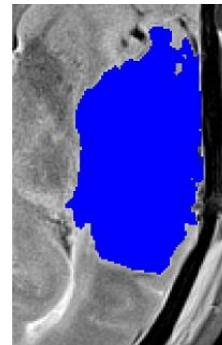
Gold standard



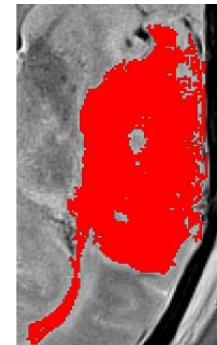
Auto Segmentation



MR image



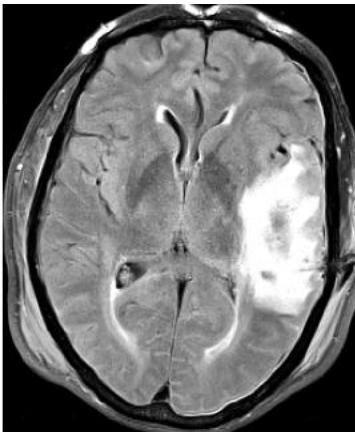
Gold standard



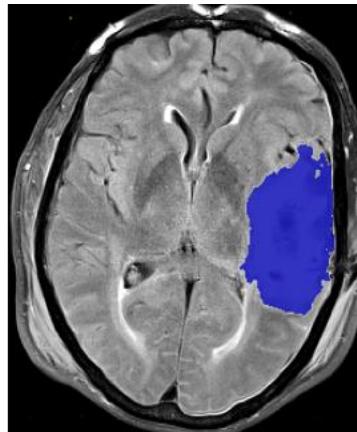
Auto Segmentation

# How to Assess Performance?

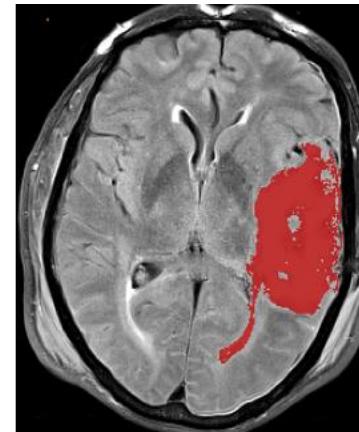
- Assume we have an automated segmentation algorithm, how do we assess "how good it is"?
- There are several ways of quantifying segmentation performance



MR image



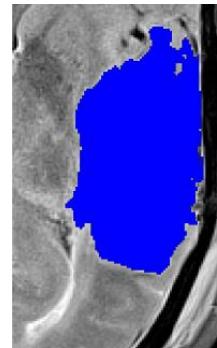
Gold standard



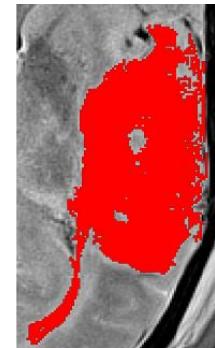
Auto Segmentation



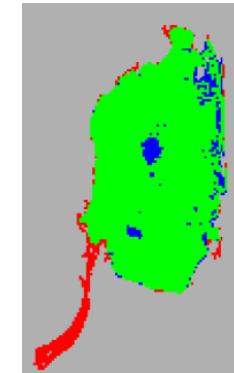
MR image



Gold standard



Auto Segmentation



TP

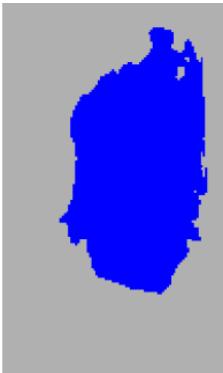
FP

FN

TN

P

N



# Confusion Matrix

---

**condition positive ()**

the number of real positive cases in the data

**condition negatives ()**

the number of real negative cases in the data

**true positive ()**

eqv. with hit

**true negative ()**

eqv. with correct rejection

**false positive ()**

eqv. with false alarm, Type I error

**false negative ()**

eqv. with miss, Type II error

		True condition	
		Total population	Condition positive
Predicted condition	Predicted condition positive	True positive	False positive, Type I error
	Predicted condition negative	False negative, Type II error	True negative

# Accuracy, Precision, Recall, ...

## accuracy

$$ACC = \frac{TP+TN}{P+N}, P = TP + FN, N = TN + FP$$

## precision or positive predictive value

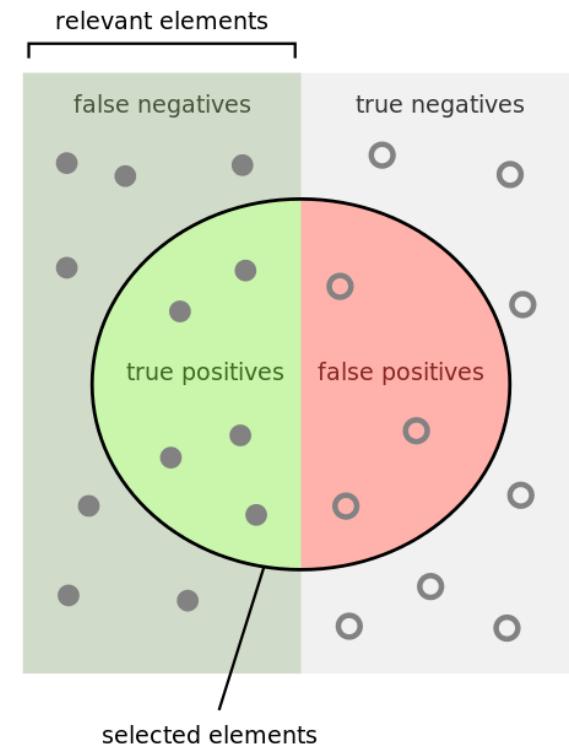
$$PPV = \frac{TP}{TP + FP}$$

## recall, sensitivity, hit rate or true positive rate

$$TPR = \frac{TP}{P} = \frac{TP}{TP + FN}$$

## specificity or true negative rate

$$TNR = \frac{TN}{N} = \frac{TN}{TN + FP}$$



How many selected items are relevant?

$$\text{Precision} = \frac{\text{green}}{\text{green} + \text{red}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{green}}{\text{green} + \text{grey}}$$

# F1 score

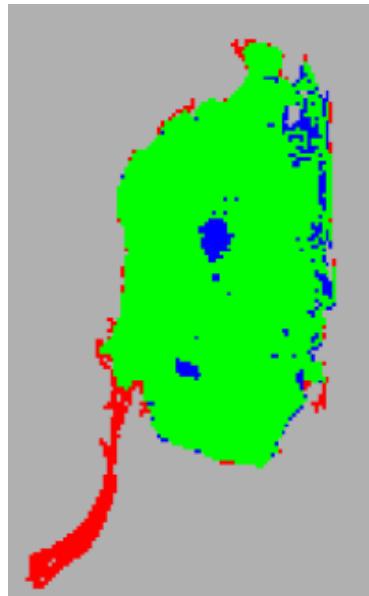
## F1 score

is the [harmonic mean](#) of precision and recall

$$F_1 = 2 \frac{PPV \cdot TPR}{PPV + TPR} = \frac{2TP}{2TP + FP + FN}$$

**precision**  $PPV = \frac{TP}{TP+FP}$

**recall**  $TPR = \frac{TP}{TP+FN}$



■ TP

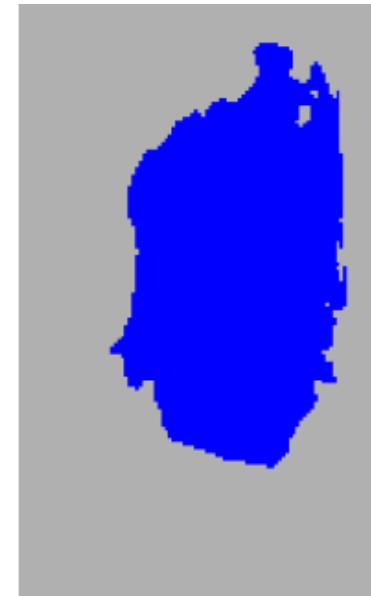
■ FP

■ FN

■ TN

P ■

N ■



# Dice Similarity Coefficient

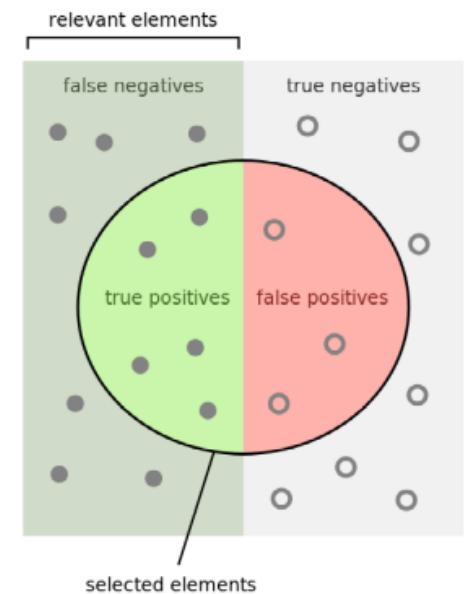
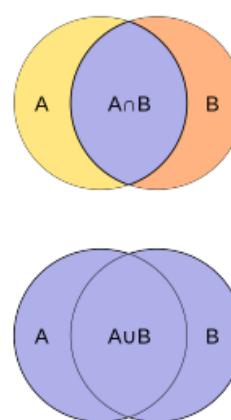
- Most widely used measure for evaluating segmentation
- Assume is the reference, and the prediction

$$DSC = \frac{2|A \cap B|}{|A| + |B|}$$

- $|A| = TP + FN$
- $|B| = TP + FP$
- $|A \cap B| = TP$

$$DSC = \frac{2TP}{2TP + FP + FN} = F_1$$

DSC is equivalent to F1 score!



# Other Measures

---

Tana & Hanbury: "Metrics for evaluating 3D medical image segmentation: analysis, selection, and tool"

## volume similarity

$$VS = 1 - \frac{||A| - |B||}{|A| + |B|} = 1 - \frac{|FN - FP|}{2TP + FP + FN}$$

## surface distance measures

- Hausdorff distance

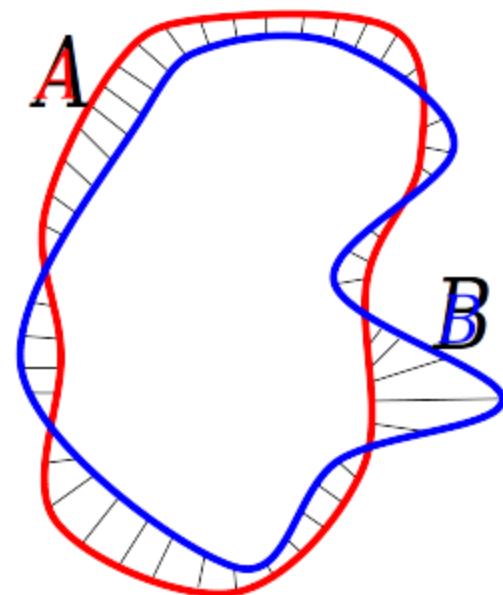
$$HD = \max(h(A, B), h(B, A))$$

$$h(A, B) = \max_{a \in A} \min_{b \in B} \|a - b\|$$

- (symmetric) average surface distance

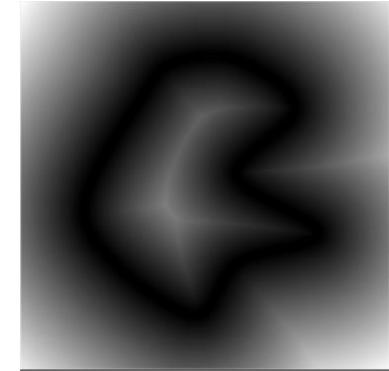
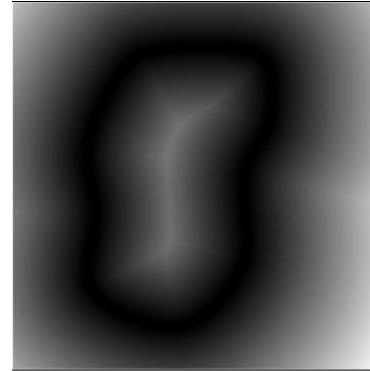
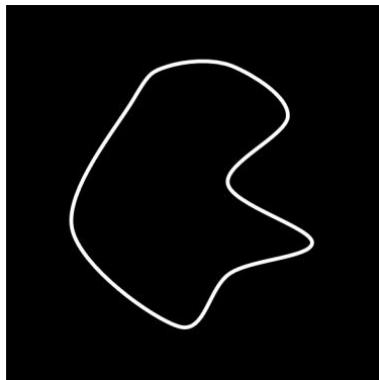
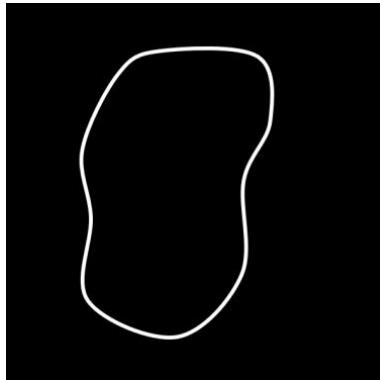
$$ASD = \frac{d(A, B) + d(B, A)}{2}$$

$$d(A, B) = \frac{1}{N} \sum_{a \in A} \min_{b \in B} \|a - b\|$$



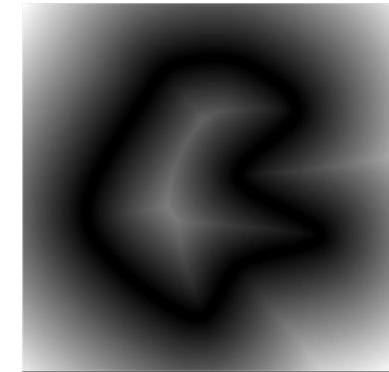
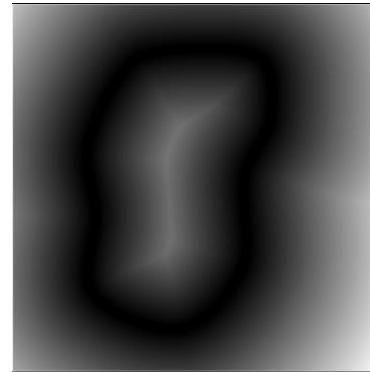
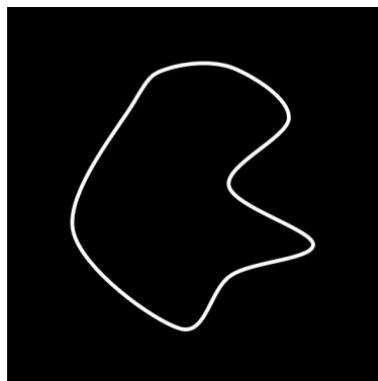
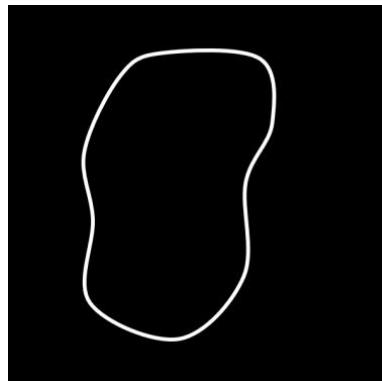
# Surface Distance

---

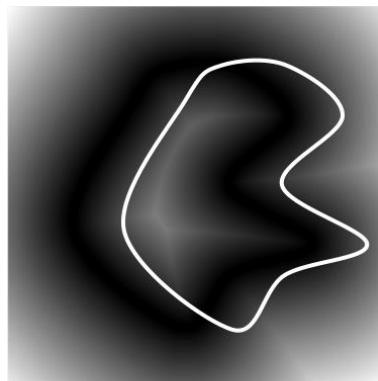
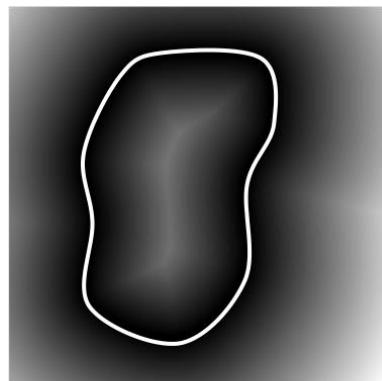


Euclidean Distance Maps

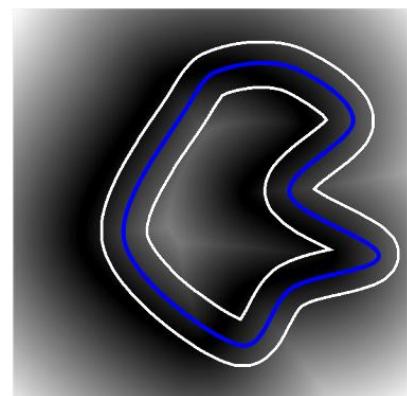
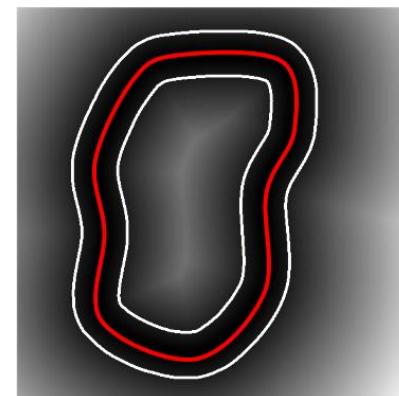
# Surface Distance



Euclidean Distance Maps

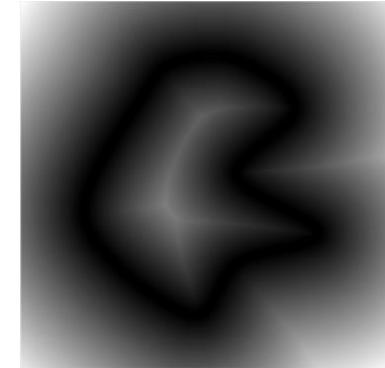
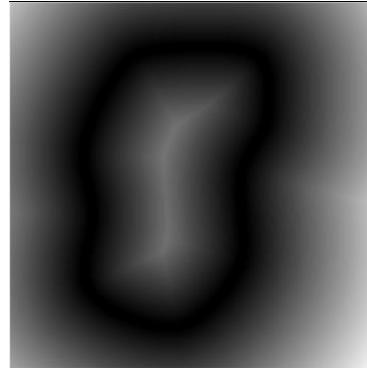
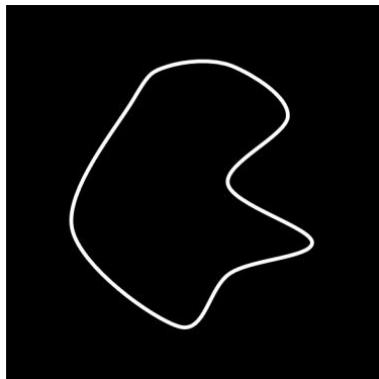
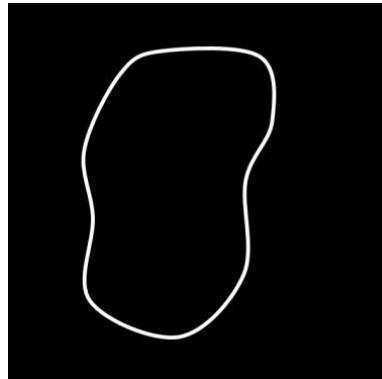


Zero-Level Surface

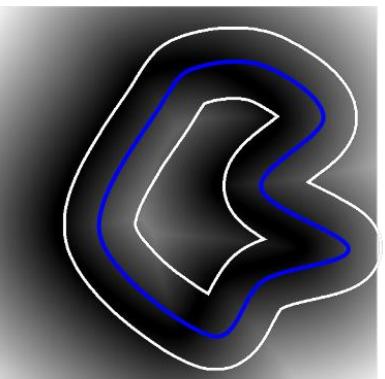
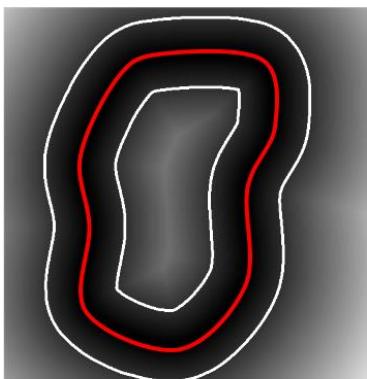
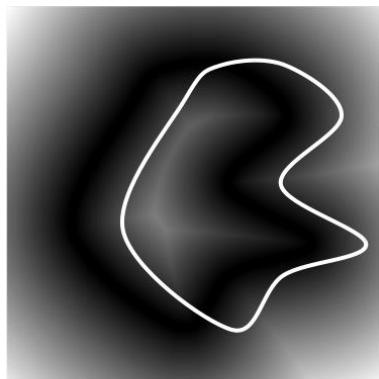
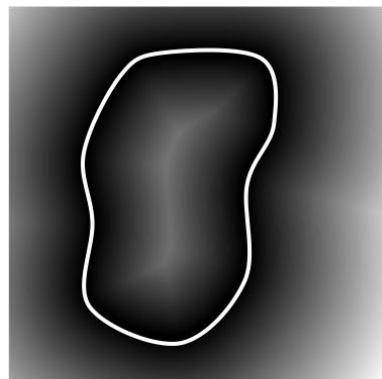


30-Pixels Distance Surface

# Surface Distance



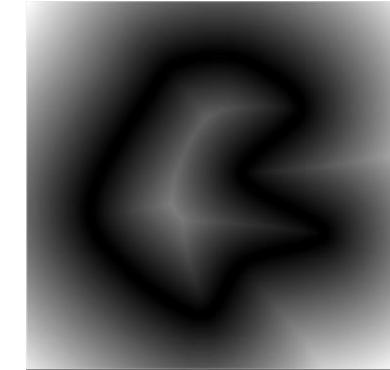
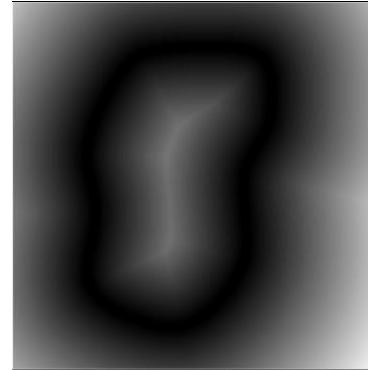
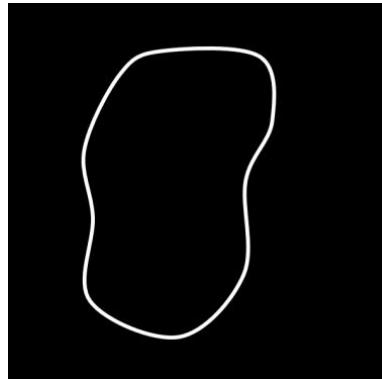
Euclidean Distance Maps



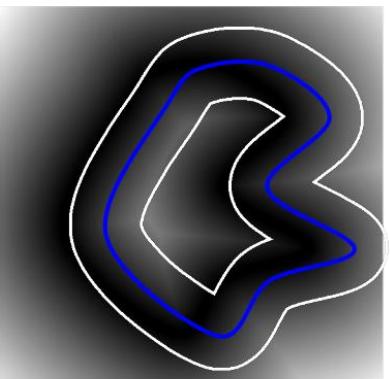
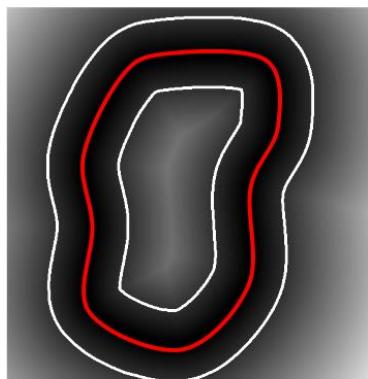
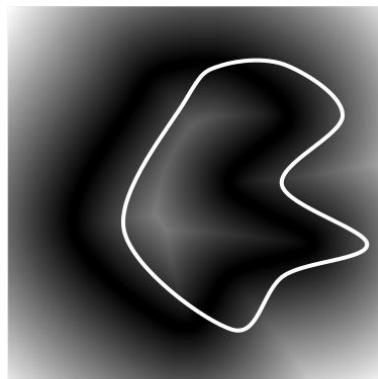
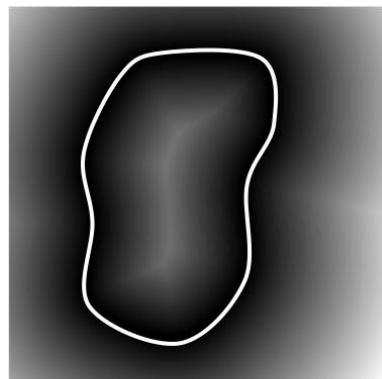
Zero-Level Surface

50-Pixels Distance Surface

# Surface Distance



Euclidean Distance Maps



Zero-Level Surface

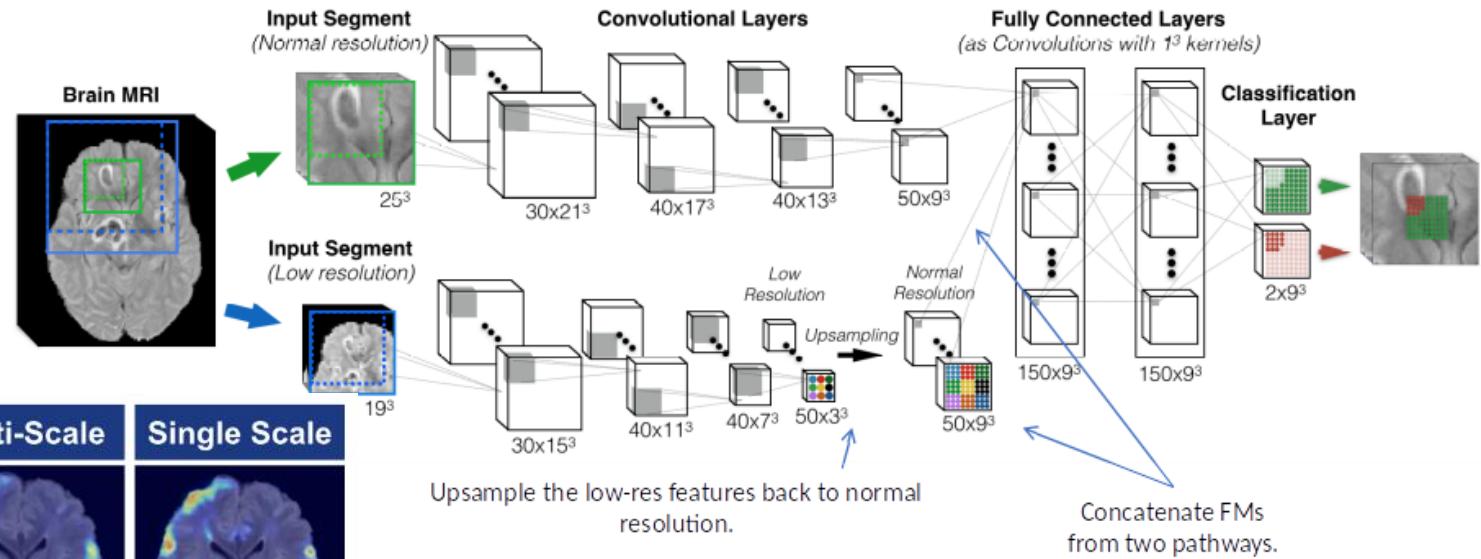
50-Pixels Distance Surface

Sum up distances along pixels on the boundaries of one contour overlaid on the distance map of the other.

# More Papers

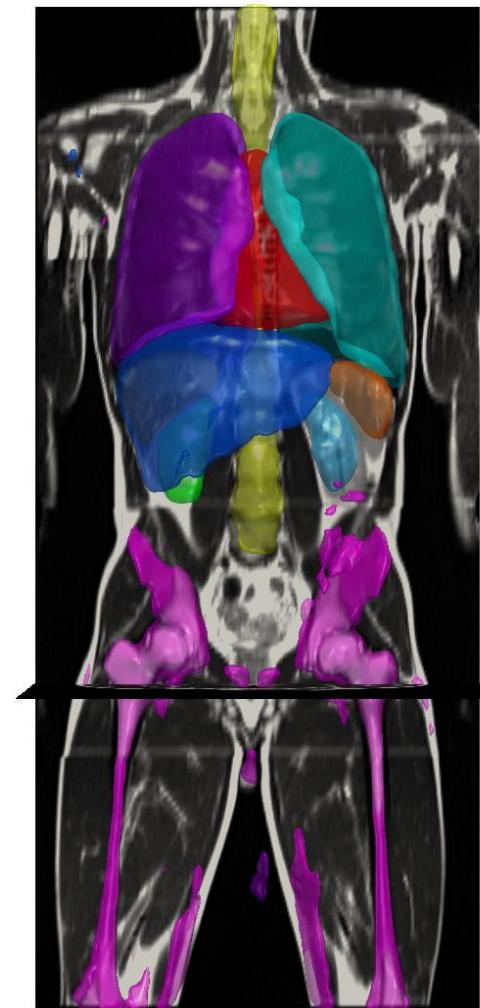
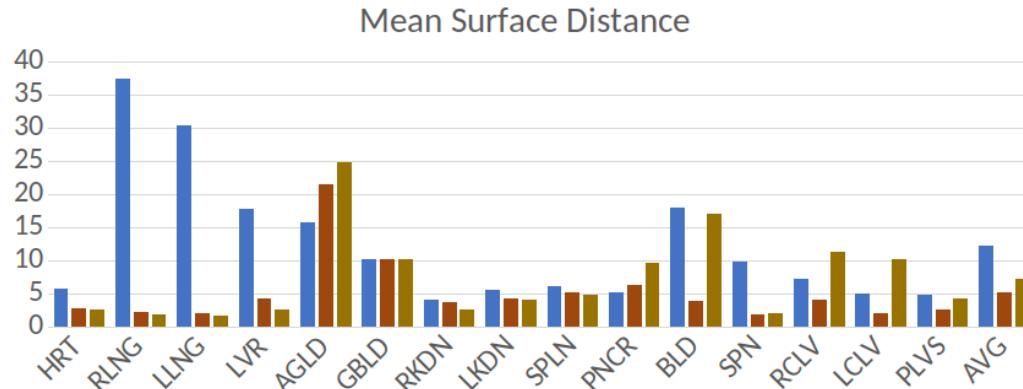
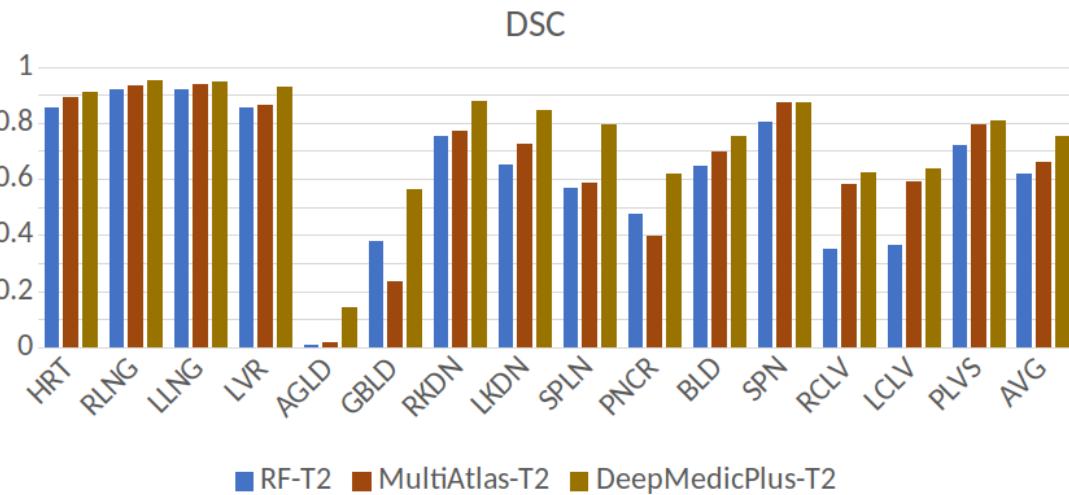
# Multi-Scale Processing

- How can we make the network to "see" more context?
- Idea: Add more pathways which process downsampled images



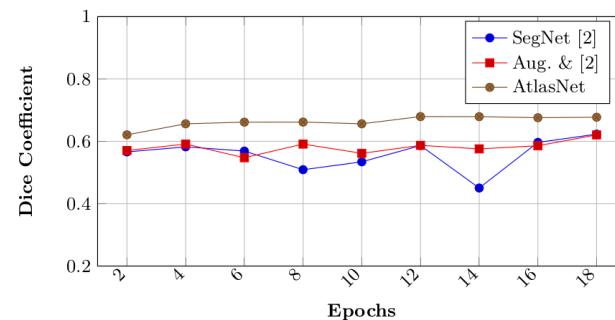
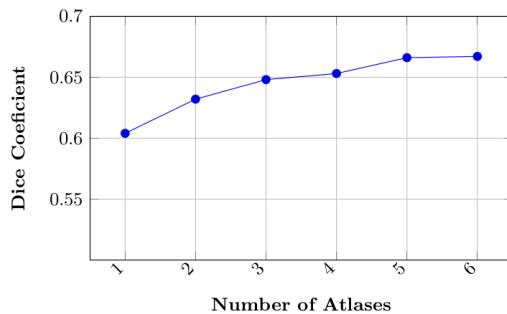
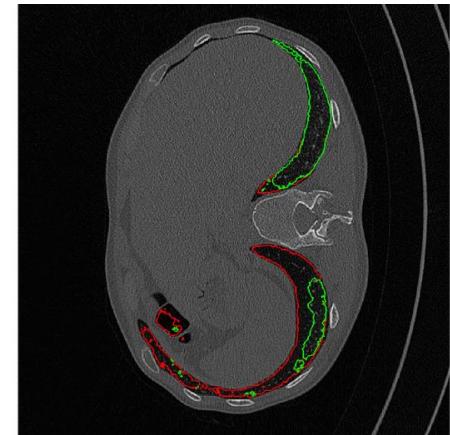
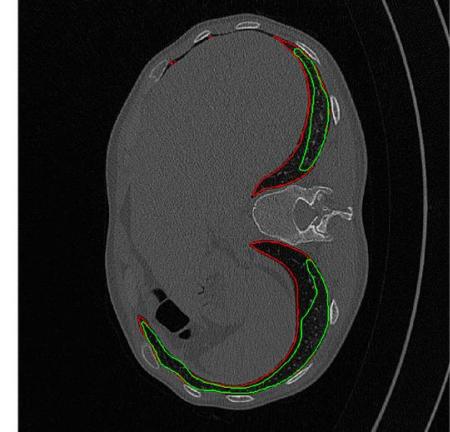
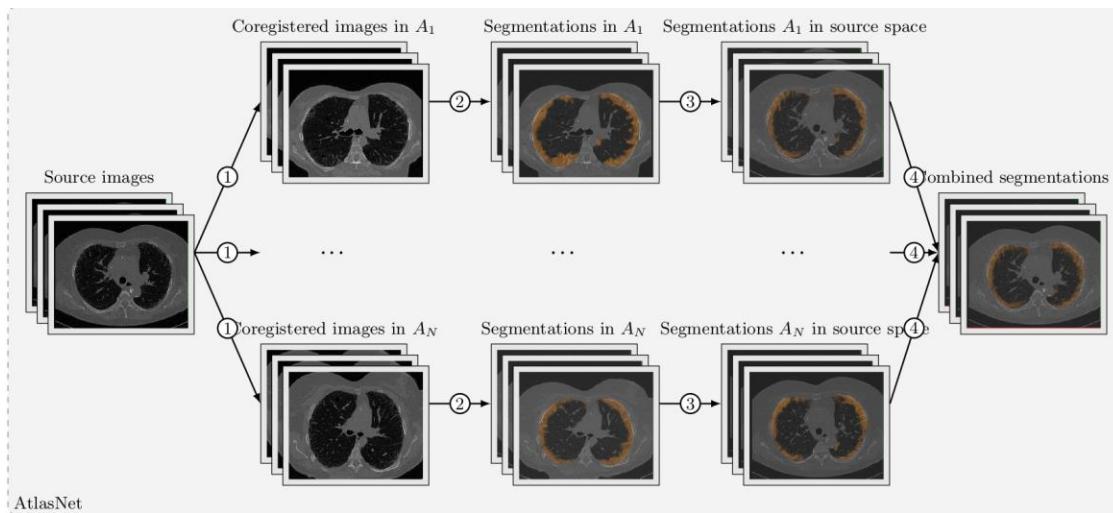
# Multi-Organ Segmentation

Lavdas et al. "Fully automatic, multiorgan segmentation in normal whole body magnetic resonance imaging (MRI), using classification forests (CFs), convolutional neural networks (CNNs), and a multi-atlas (MA) approach"



# AtlasNet

Vakalopoulou et al. "AtlasNet: Multi-atlas Non-linear Deep Networks for Medical Image Segmentation"



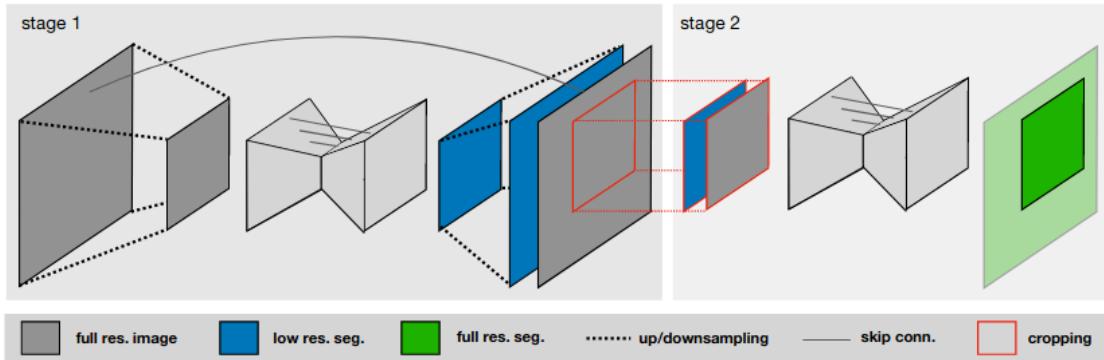
# nnU-Net

Isensee et al. "nnU-Net: Self-adapting Framework for U-Net-Based Medical Segmentation"

$$\mathcal{L}_{total} = \mathcal{L}_{dice} + \mathcal{L}_{CE}$$

		2D U-Net	3D U-Net	3D U-Net lowres	
BrainTumour	median patient shape	169x138	138x169x138	-	
	input patch size	192x160	128x128x128	-	
	batch size	89	2	-	
	num pool per axis	5, 5	5, 5, 5	-	
Heart	median patient shape	320x232	115x320x232	58x160x116	
	input patch size	320x256	80x192x128	64x160x128	
	batch size	33	2	2	
	num pool per axis	6, 6	4, 5, 5	4, 5, 5	
Liver	median patient shape	512x512	482x512x512	121x128x128	
	input patch size	512x512	128x128x128	128x128x128	
	batch size	10	2	2	
	num pool per axis	6, 6	5, 5, 5	5, 5, 5	
Hippocampus	median patient shape	50x35	36x50x35	-	
	input patch size	56x40	40x56x40	-	
	batch size	366	9	-	
	num pool per axis	3, 3	3, 3, 3	-	
Prostate	median patient shape	320x319	20x320x319	-	
	input patch size	320x320	20x192x192	-	
	batch size	26	4	-	
	num pool per axis	6, 6	2, 5, 5	-	
Lung	median patient shape	512x512	252x512x512	126x256x256	
	input patch size	512x512	112x128x128	112x128x128	
	batch size	10	2	2	
	num pool per axis	6, 6	4, 5, 5	4, 5, 5	
Pancreas	median patient shape	512x512	96x512x512	96x256x256	
	input patch size	512x512	96x160x128	96x160x128	
	batch size	10	2	2	
	num pool per axis	6, 6	4, 5, 5	4, 5, 5	

label	BrainTumour		Heart	Liver	Hippoc.	Prostate	Lung	Pancreas				
	1	2	3	1	1	2	1	2	1	2	1	2
2D U-Net	78.60	58.65	77.42	91.36	94.37	53.94	88.52	86.70	61.98	84.31	52.68	74.70
3D U-Net	<b>80.71</b>	<b>62.22</b>	<b>79.07</b>	92.45	94.11	61.74	<b>89.87</b>	<b>88.20</b>	60.77	83.73	55.87	77.69
3D U-Net stage1 only (U-Net Cascade)	-	-	-	90.63	94.69	47.01	-	-	-	-	65.33	79.45
3D U-Net (U-Net Cascade)	-	-	-	92.40	95.38	58.49	-	-	-	-	<b>66.85</b>	<b>79.30</b>
ensemble												
2D U-Net+	80.79	61.72	79.16	<b>92.70</b>	94.30	60.24	89.78	88.09	<b>63.78</b>	<b>85.31</b>	55.96	78.26
3D U-Net												
ensemble												
2D U-Net+	-	-	-	92.64	95.31	60.09	-	-	-	-	61.18	78.79
3D U-Net (U-Net Cascade)	-	-	-				-	-	-	-	65.16	79.70
ensemble												
3D U-Net+	-	-	-	92.63	<b>95.43</b>	<b>61.82</b>	-	-	-	-	65.16	79.70
3D U-Net (U-Net Cascade)	-	-	-				-	-	-	-	49.14	
test set	67.71	<b>47.73</b>	<b>68.16</b>	<b>92.77</b>	<b>95.24</b>	<b>73.71</b>	<b>90.37</b>	<b>88.95</b>	<b>75.81</b>	<b>89.59</b>	<b>69.20</b>	<b>79.53</b>



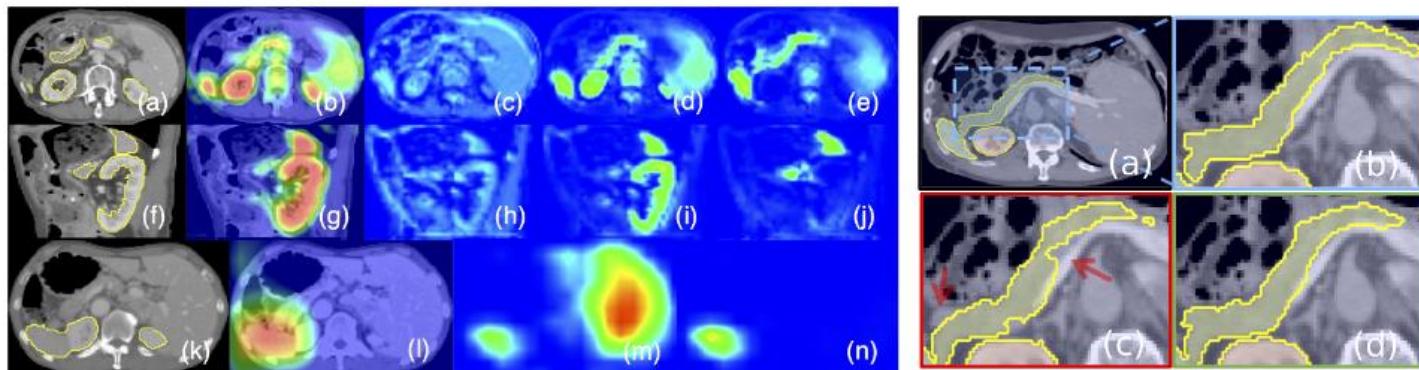
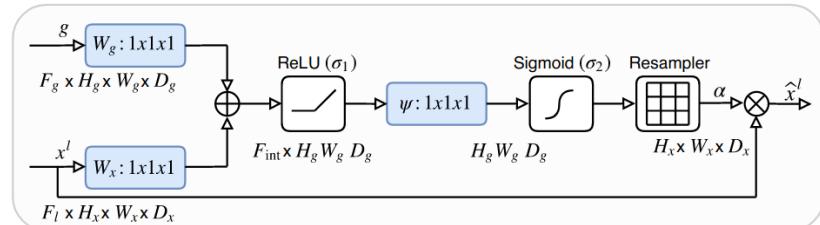
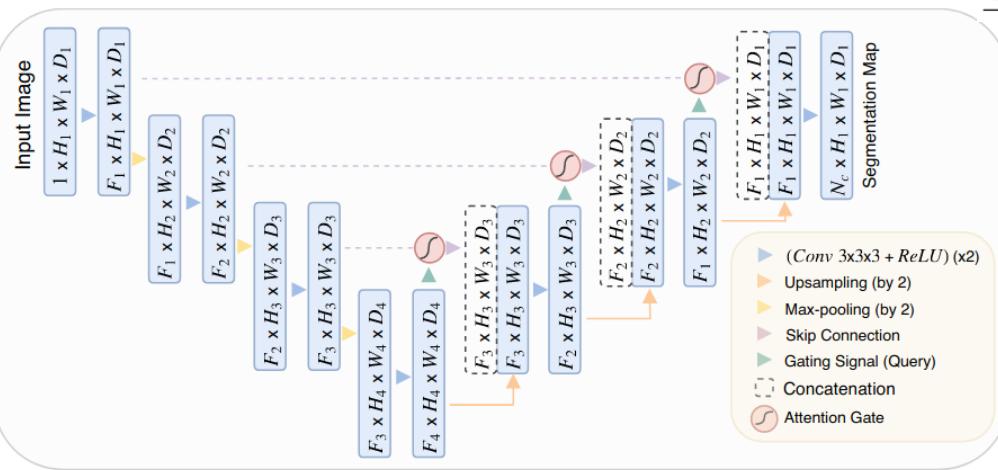
# Attention U-Net

Oktay et al. "Attention U-Net: Learning Where to Look for the Pancreas"

$$q_{att}^l = \psi^T (\sigma_1 (W_x^T x_i^l + W_g^T g_i + b_g)) + b_\psi$$

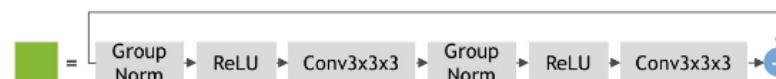
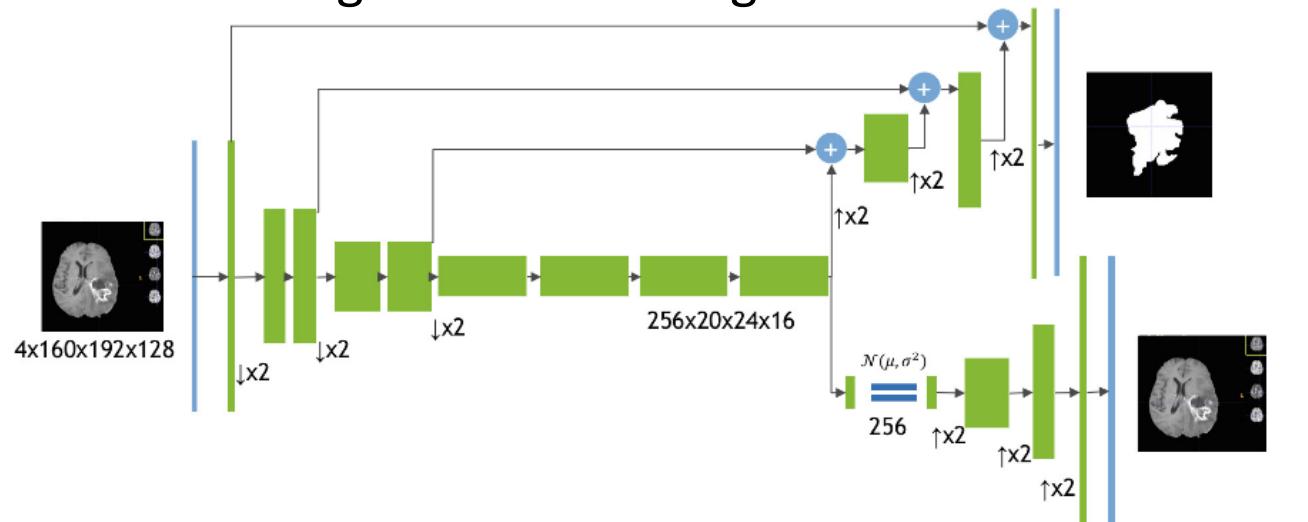
$$\alpha_i^l = \sigma_2(q_{att}^l(x_i^l, g_i; \Theta_{att})),$$

	Method	Dice Score	Precision	Recall	S2S Dist (mm)
BFT	U-Net [24]	0.690±0.132	0.680±0.109	0.733±0.190	6.389±3.900
	Attention U-Net	<b>0.712±0.110</b>	0.693±0.115	<b>0.751±0.149</b>	<b>5.251±2.551</b>
AFT	U-Net [24]	0.820±0.043	0.824±0.070	0.828±0.064	2.464±0.529
	Attention U-Net	<b>0.831±0.038</b>	0.825±0.073	<b>0.840±0.053</b>	<b>2.305±0.568</b>
SCR	U-Net [24]	0.815±0.068	0.815±0.105	0.826±0.062	2.576±1.180
	Attention U-Net	0.821±0.057	0.815±0.093	<b>0.835±0.057</b>	<b>2.333±0.856</b>



# Multi-Task

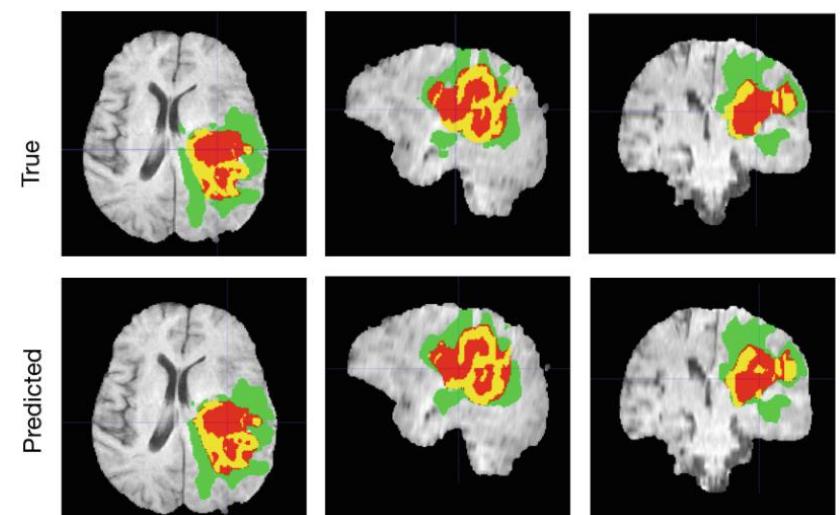
Myronenko "3D MRI Brain Tumor Segmentation Using Autoencoder Regularization"



$\downarrow \times 2 = \text{conv}3 \times 3 \times 3 \text{ stride } 2$

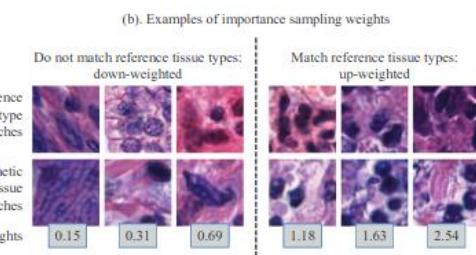
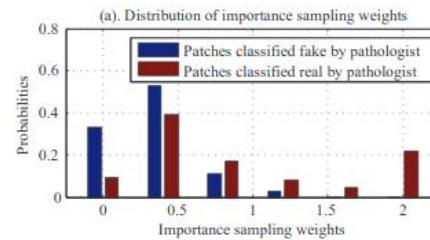
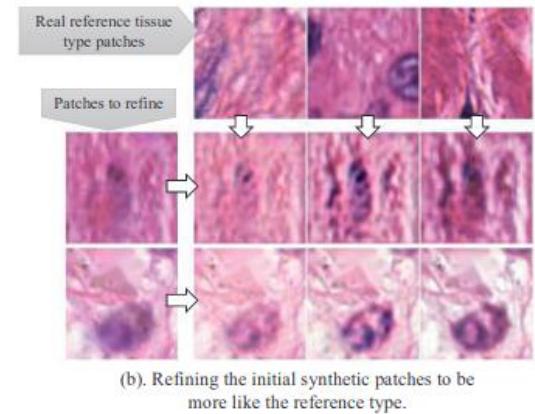
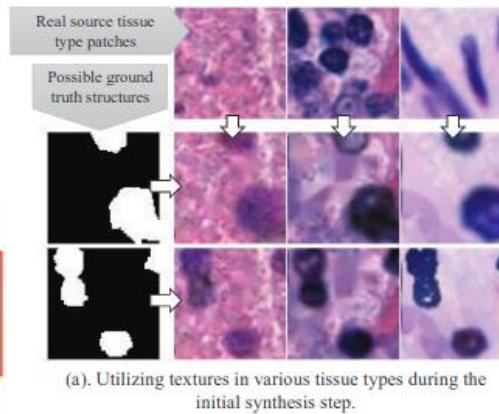
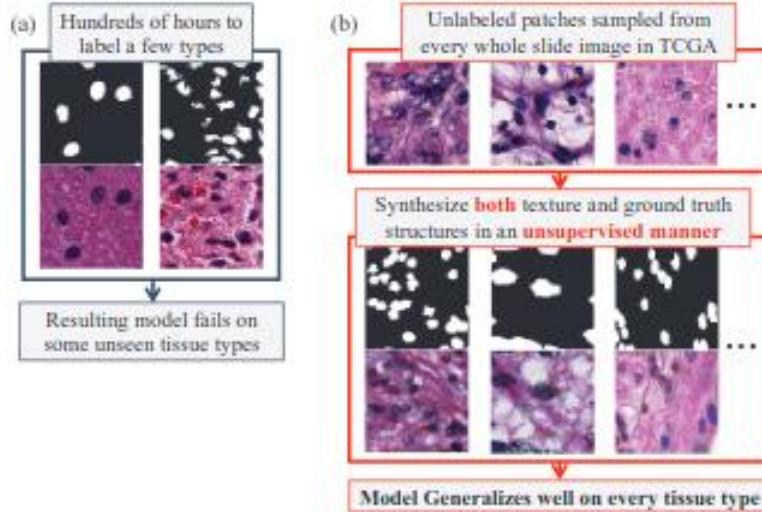
$\uparrow \times 2 = \text{conv}1 \times 1 \times 1, 3\text{D bilinear upsizing}$

$$\mathbf{L} = \mathbf{L}_{dice} + 0.1 * \mathbf{L}_{L2} + 0.1 * \mathbf{L}_{KL}$$



# Synthesize

Hou et al. "Robust Histopathology Image Analysis: to Label or to Synthesize?"

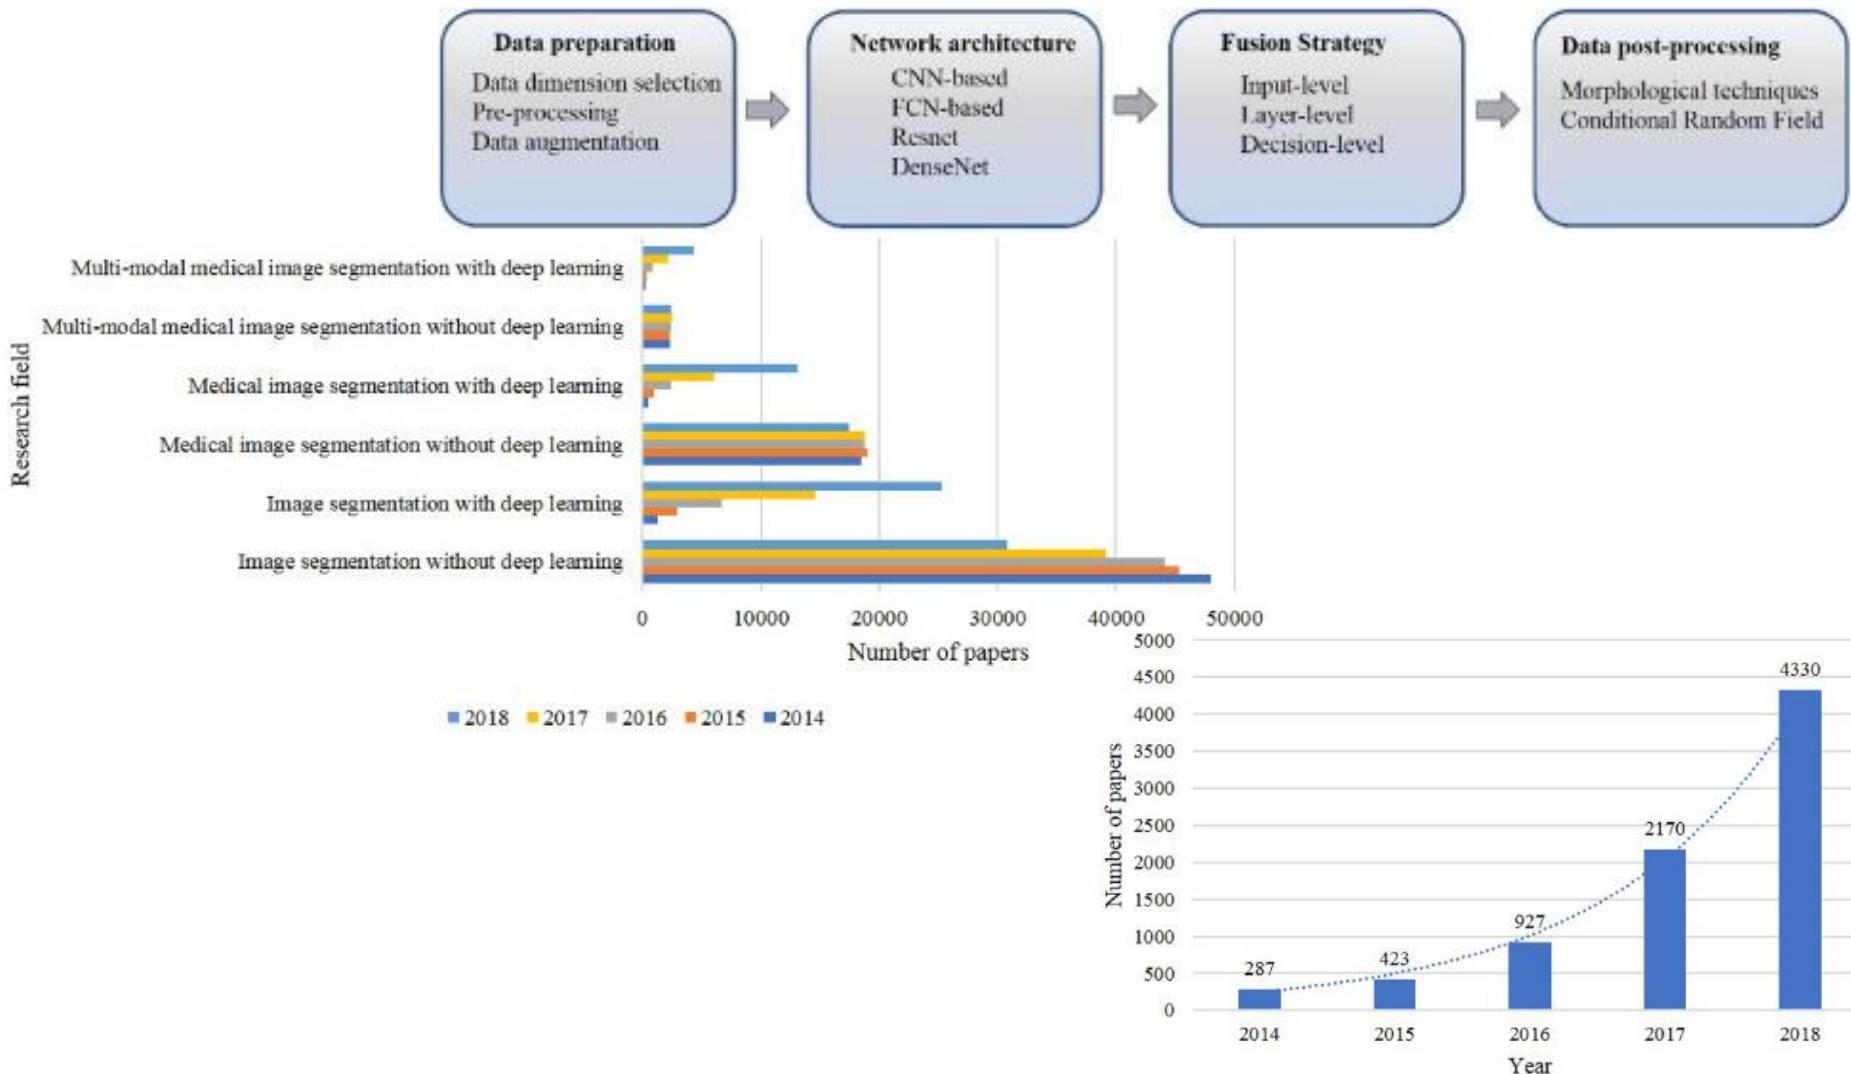


Nucleus segmentation methods	DICE Avg.
No hard examples	0.7476
No reference patch during refinement	0.7410
No importance weights	0.7533
Universal CNN (proposed)	<b>0.7612</b>

Lymphocyte detection methods	AUROC
Level Set features + supervised net [67]	0.7132
Fine-tuning VGG16 (supervised) [52]	0.6925
Universal CNN (proposed)	<b>0.7149</b>

# Interesting Points

Zhou et al. "A review: Deep Learning for medical image segmentation using multi-modality fusion"



# Agony of Choice

---

- Architecture: depth, width, scales, residuals, ....
- Loss function: (weighted) cross entropy, IoU, Dice, ...
- Sampling strategy: equally per class, fore/background, uniform, ...
- Optimization: optimizer, learning rate, momentum, regularization, ...
- Data normalization: z-score, bias field correction, histogram matching, ...
- Post-processing: CRFs, ensemble of networks, ....

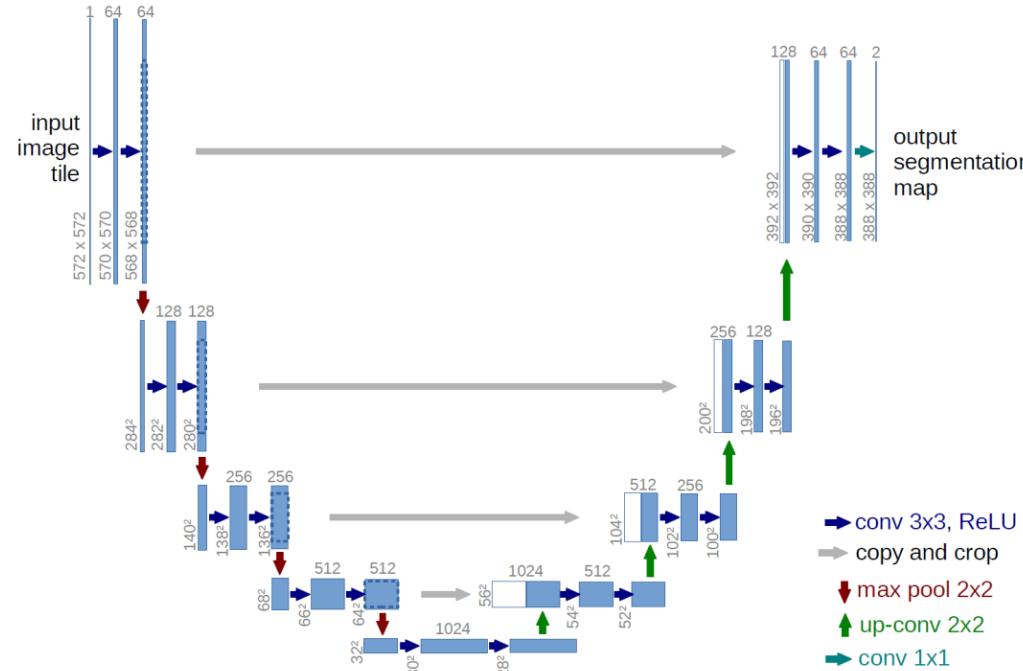
# Conclusions

---

- Semantic Segmentation is one of the most studied problems in medical imaging.
  - Conducting **quantitative analyses**, e.g. measuring the volume of the ventricular cavity.
  - Determining the precise **location and extent** of an organ or a certain type of tissue, e.g. a tumour for treatment such as radiation therapy.
  - Creating **3D models** used for **simulation**, e.g. generating a model of an abdominal aortic aneurysm for simulating stress

# Lab Session!

- Brain tumor segmentation
  - U-Net

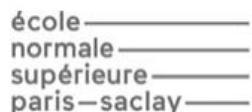


TorchVision

# Deep learning for medical imaging

**Olivier Colliot, PhD**  
**Research Director at CNRS**  
Co-Head of the ARAMIS Lab –  
[www.aramislab.fr](http://www.aramislab.fr)  
PRAIRIE – Paris Artificial Intelligence  
Research Institute

**Maria Vakalopoulou, PhD**  
**Assistant Professor at**  
**CentraleSupélec**  
Mathematics and Informatics (MICS)  
Office: Bouygues Building Sb.132



## Master 2 - MVA

Course website: <http://www.aramislab.fr/teaching/DLMI-2019-2020/>  
Piazza (for registered students):  
<https://piazza.com/centralesupelec/spring2020/mvadlmi/>